

A man with a beard and short dark hair is wearing a white and black VR headset. He is gesturing with both hands raised in front of him, palms facing each other. In the foreground, there is a detailed wooden architectural model of a multi-story building with various rooms and windows. The background is a blurred indoor setting with shelves and other objects.

arm

Arm and SVE

International Workshop on High-
Performance Computing and
Programming on Quantum Chemistry
and Physics 2020 (HPCPQCP2020)

Dr. Olly Perks (olly.perks@arm.com)
15th January 2020

Agenda

- Part 1: Getting to SVE (45 Mins)
 - A quick introduction of Arm in HPC / Scientific Computing
 - Existing Arm vector units / instructions: NEON
 - The SVE specification
- Part 2: Going Forward with SVE (45 mins)
 - Using SVE for real applications
 - Next generation vector instructions: Beyond SVE

A woman with long blonde hair is shown in profile, looking towards the right. She is standing in front of a large digital display or window. The entire image is overlaid with a teal/cyan color filter. In the background, there are blurred lights, suggesting an urban or office setting at night. The woman's hand is visible, touching the surface of the display. The text 'arm' is on the left, and 'Part 1: Getting to SVE' is on the right.

arm

Part 1: Getting to SVE

arm

Recap:
Arm and Arm in HPC

Arm Technology Already Connects the World

An aerial night view of New York City, showing the dense urban landscape with numerous illuminated skyscrapers and buildings. Overlaid on the city is a network diagram consisting of white dots connected by thin white lines, forming a web that spans across the city and extends towards the horizon, symbolizing global connectivity.

Arm is ubiquitous

21 billion chips sold by partners in 2017 alone

Mobile/Embedded/IoT/
Automotive/Server/GPUs

Partnership is key

We design IP, not
manufacture chips

Partners build products
for their target markets

Choice is good

One size is not always the
best fit for all

HPC is a great fit for
co-design and collaboration

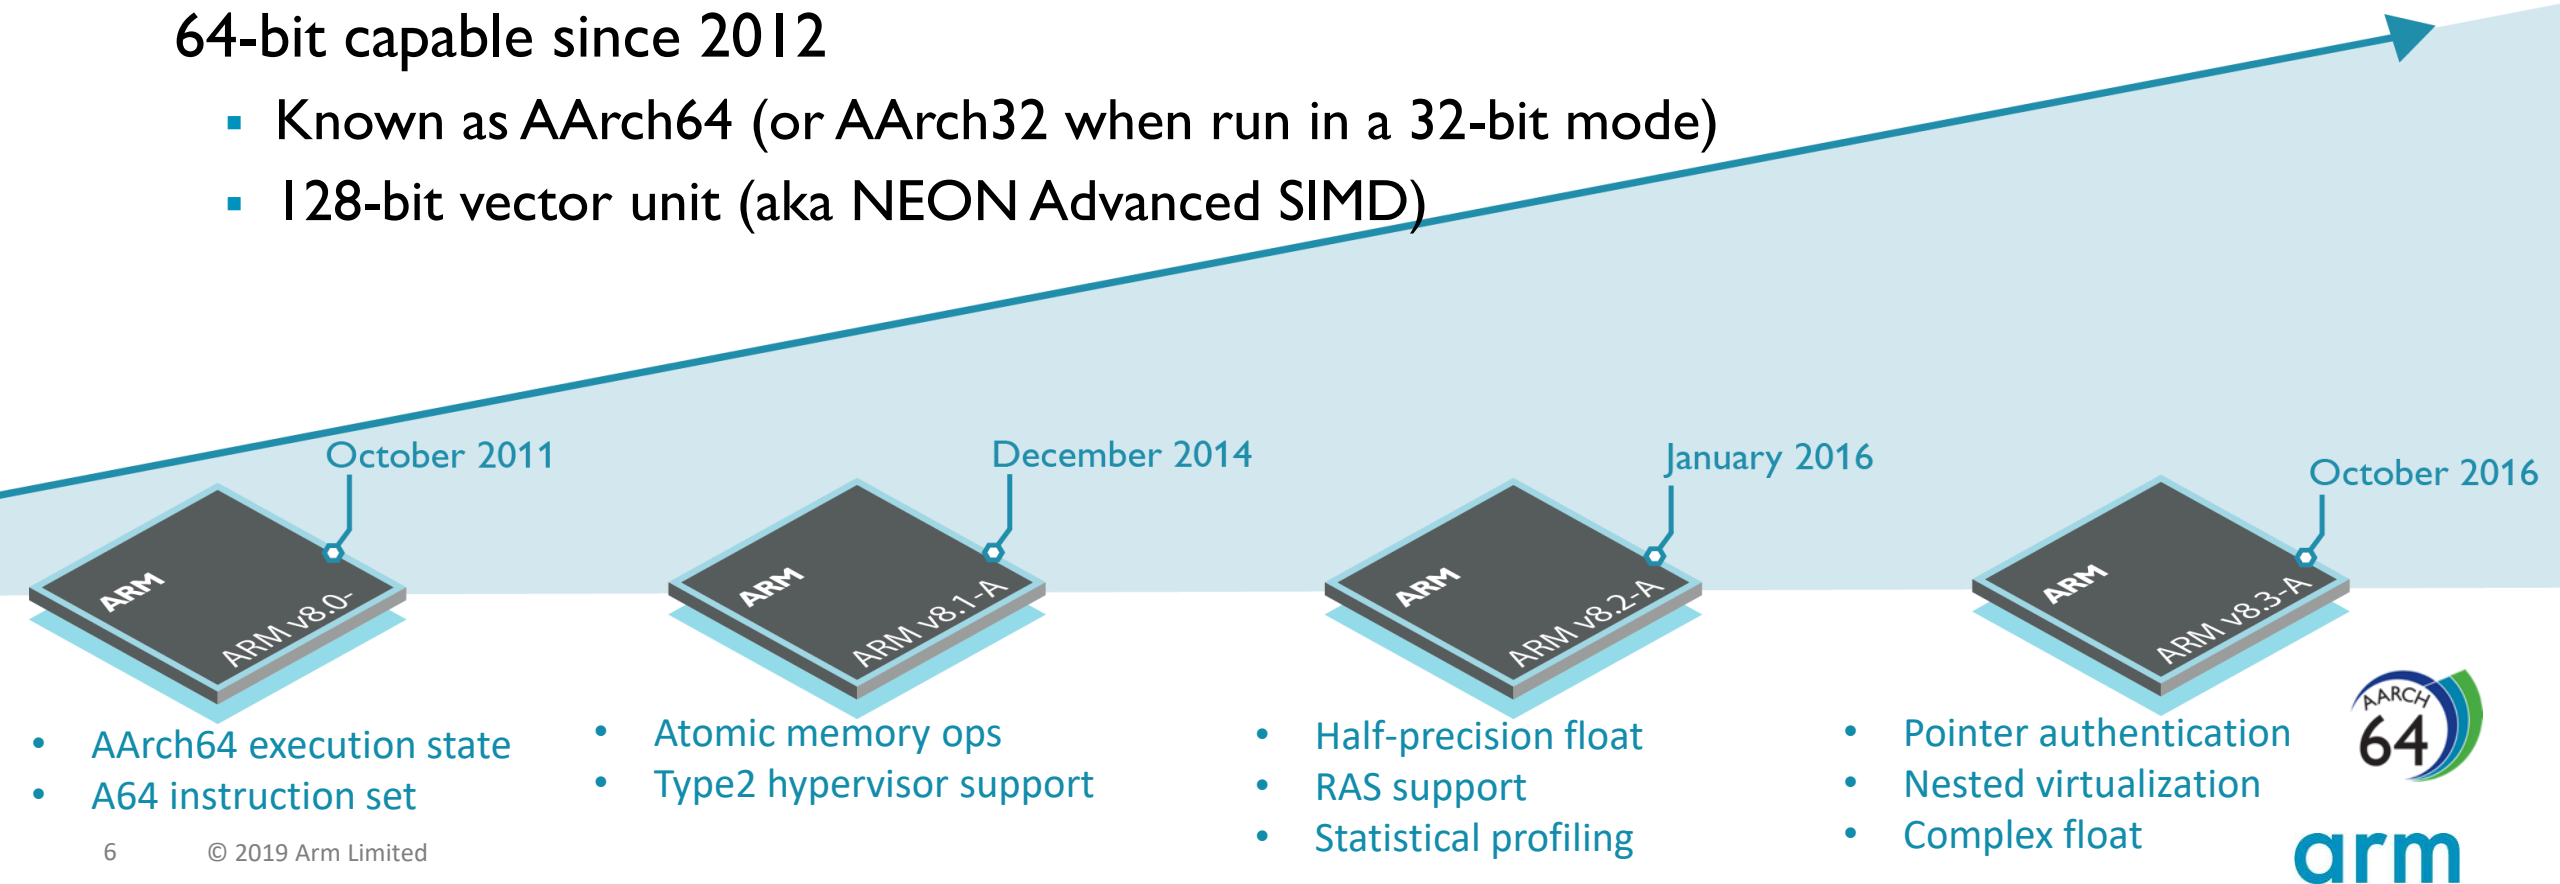
Armv8-A Architecture Evolution

RISC architecture


- Only have 32 bits available for encoding all instructions
- Supports the development of efficient implementations

64-bit capable since 2012

- Known as AArch64 (or AArch32 when run in a 32-bit mode)
- 128-bit vector unit (aka NEON Advanced SIMD)



History of Arm in HPC: A Busy Decade



2011 Calxeda

- 32-bit ARmv7-A – Cortex A9



2011-2015
Mont-Blanc 1

- 32-bit Armv7-A
- Cortex A15
- First Arm HPC system




2014
AMD Opteron
A1100

- 64-bit Armv8-A
- Cortex A57
- 4-8 Cores




2015
Cavium
ThunderX

- 64-bit Armv8-A
- 48 Cores




2017
(Cavium)
Marvell
ThunderX 2

- 64-bit Armv8-A
- 32 Cores



History of Arm in HPC: A Busy Decade



2011 Calxeda

- 32-bit ARmv7-A – Cortex A9



2011-2015
Mont-Blanc 1

- 32-bit Armv7-A
- Cortex A15
- First Arm HPC system




2014
AMD Opteron
A1100

- 64-bit Armv8-A
- Cortex A57
- 4-8 Cores



2015
Cavium
ThunderX

- 64-bit Armv8-A
- 48 Cores



2017
(Cavium)
Marvell
ThunderX 2

- 64-bit Armv8-A
- 32 Cores



2019
Fujitsu A64FX

- First Arm chip with SVE vectorisation
- 48 Cores



Why Arm?

Especially for Infrastructure / HPC / Scientific Computing / ML?

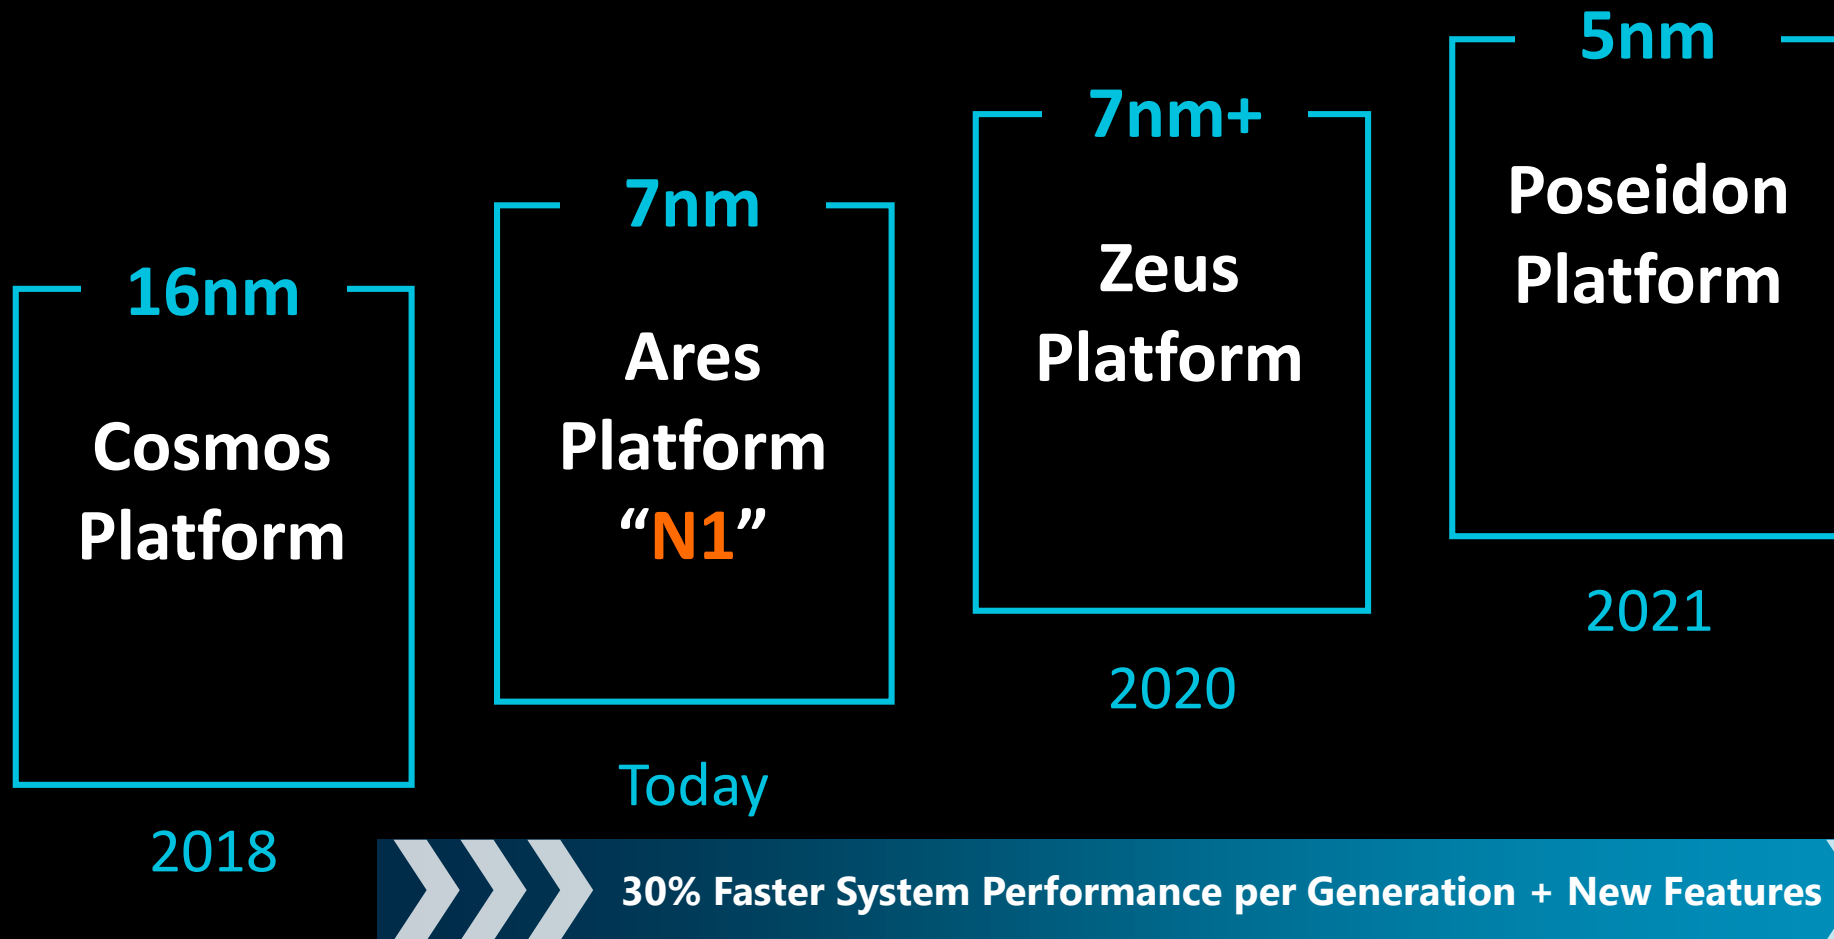
Hardware

- Flexibility: Allow vendors to differentiate
 - Speed and cost of development
- Provide different licensing
 - Core - Reference design (A53/A72/N1)
 - Architecture - Design your own (TX2, A64FX)
- Other hardware components
 - NoCs, GPUs, memory controllers
 - “Building blocks” design
- Architecture validation for correctness












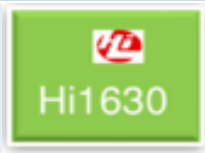





Software

- All based on the same instruction set
 - Commonality between hardware
 - Reuse of software
- Comprehensive software ecosystem
 - Operating systems, compilers, libraries, tools
 - Not just vendor - third party too
- Large community
 - Everything from Android to HPC

Each generation brings faster performance and new infrastructure specific features



Variation in the Processor Market

Marvell (Cavium)						
Ampere (X-Gene)						
Fujitsu						
Huawei (HiSilicon)						
Amazon (Annapurna)						
EPI / SiPearl						
Other						

Arm Processor Top Trumps

Marvell: ThunderX2



Core Count:	32
Clock Speed:	2.5 GHz
Memory Bandwidth:	~130 GB/s
Energy Consumption:	~200 w
Vector Units:	128-bit NEON

Huawei: Kunpeng 920



Core Count:	64
Clock Speed:	2.6 GHz
Memory Bandwidth:	~130 GB/s
Energy Consumption:	~180 w
Vector Units:	128-bit NEON

Fujitsu: A64FX



Core Count:	48
Clock Speed:	2 GHz
Memory Bandwidth:	~1 TB/s
Energy Consumption:	~150 w
Vector Units:	512-bit SVE

HPC Deployments

- More Arm based CPUs are being adopted
 - Lots of large scale deployments
- Different OEMs
 - Cray, HPE, Atos-Bull, Fujitsu, Huawei, E4
- EU Deployments
 - Isambard: Cray 10k TX2 cores
 - GENCIE: Atos 18k TX2 cores
 - Catalyst 3 systems: HPE 4k TX2 core
 - Fugaku Prototype: Fujitsu 36.8K A64FX cores
 - **Future** Isambard 2: Cray A64FX
 - **Future** Deucalion: Cray A64FX



>5k ThunderX2 CPUs



2k Kunpeng 920 CPUs + 8k AI accelerators



150k+ Fujitsu A64FX CPUs



The Cloud

Open access to server class Arm

VERNE
GLOBAL

Verne Global: What is hpcDIRECT? In partnership with Arm and

VERNE GLOBAL

in partnership

arm



HPCwire

Since 1987 - Covering the Fastest Computers
in the World and the People Who Run Them



AWS Graviton Processor

First Arm Cloud Instances

POWERED BY AWS GRAVITON2 PROCESSORS

M6g, R6g, C6g
instances for EC2

New generation of Arm-based instances powered by
AWS Graviton2 processors offer 40% better
price/performance than current x86-based instances

M6g PREVIEW TODAY R6g COMING SOON C6g COMING SOON



packet

c1.large.arm

With 96 physical Arm cores, this server is anything but a lightweight - and it comes with 128 GB of RAM for just \$0.50/hr. Nice!


arm

NAMD 3 Alpha, Tesla V100 Performance, ApoA1 92k Atoms

Hardware platform		ns/day,	Speedup
Intel Xeon 8168	+ 1x Tesla V100	102.1,	1.0x
IBM Power9	+ 1x Tesla V100	99.7,	0.97x
Cavium ThunderX2	+ 1x Tesla V100	98.2,	0.96x

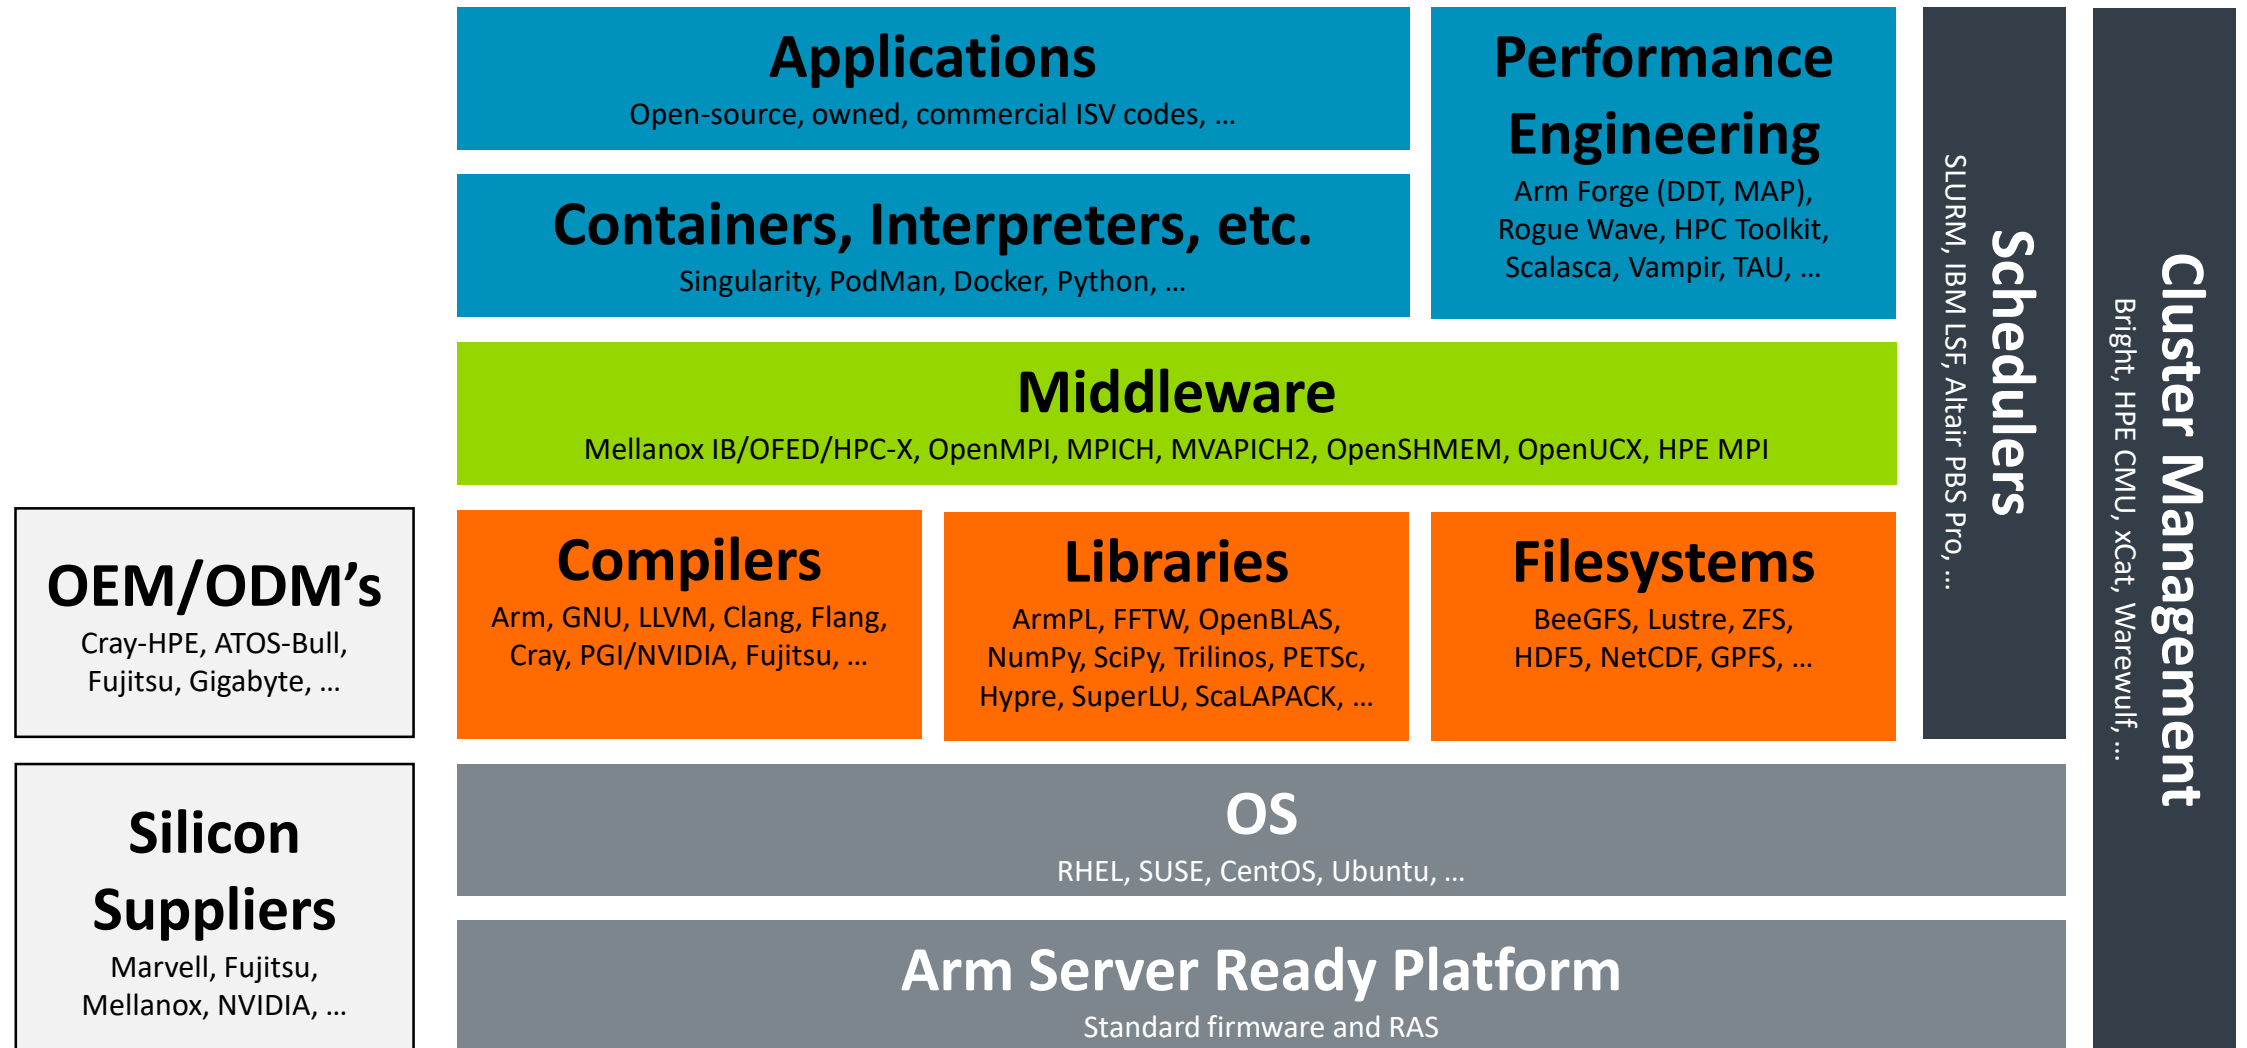
NAMD 3 Alpha Note:

This measured the new NAMD 3 Alpha in single-node mode, yielding greatest overall GPU acceleration for small to moderate sized atomic structure simulations.



ARM+NVIDIA
Preliminary Results

Not Just Hardware



A Rich and Growing Application Ecosystem


GROMACS	LAMMPS	CESM2	MrBayes	Bowtie
NAMD	AMBER	Paraview	SIESTA	UM
WRF	Quantum ESPRESSO	VASP	MILC	GEANT4
OpenFOAM	GAMESS	VisIt	DL-Poly	NEMO
BLAST	NWCHEM	Abinit	BWA	QMCPACK
Chem/Phys	Weather	CFD	Visualization	Genomics



Community driven Applications recipes:

<https://gitlab.com/arm-hpc/packages/-/wikis/categories/allPackages>

Single node performance results



Arm Provided HPC Software

Just from Arm, other vendors provide their own

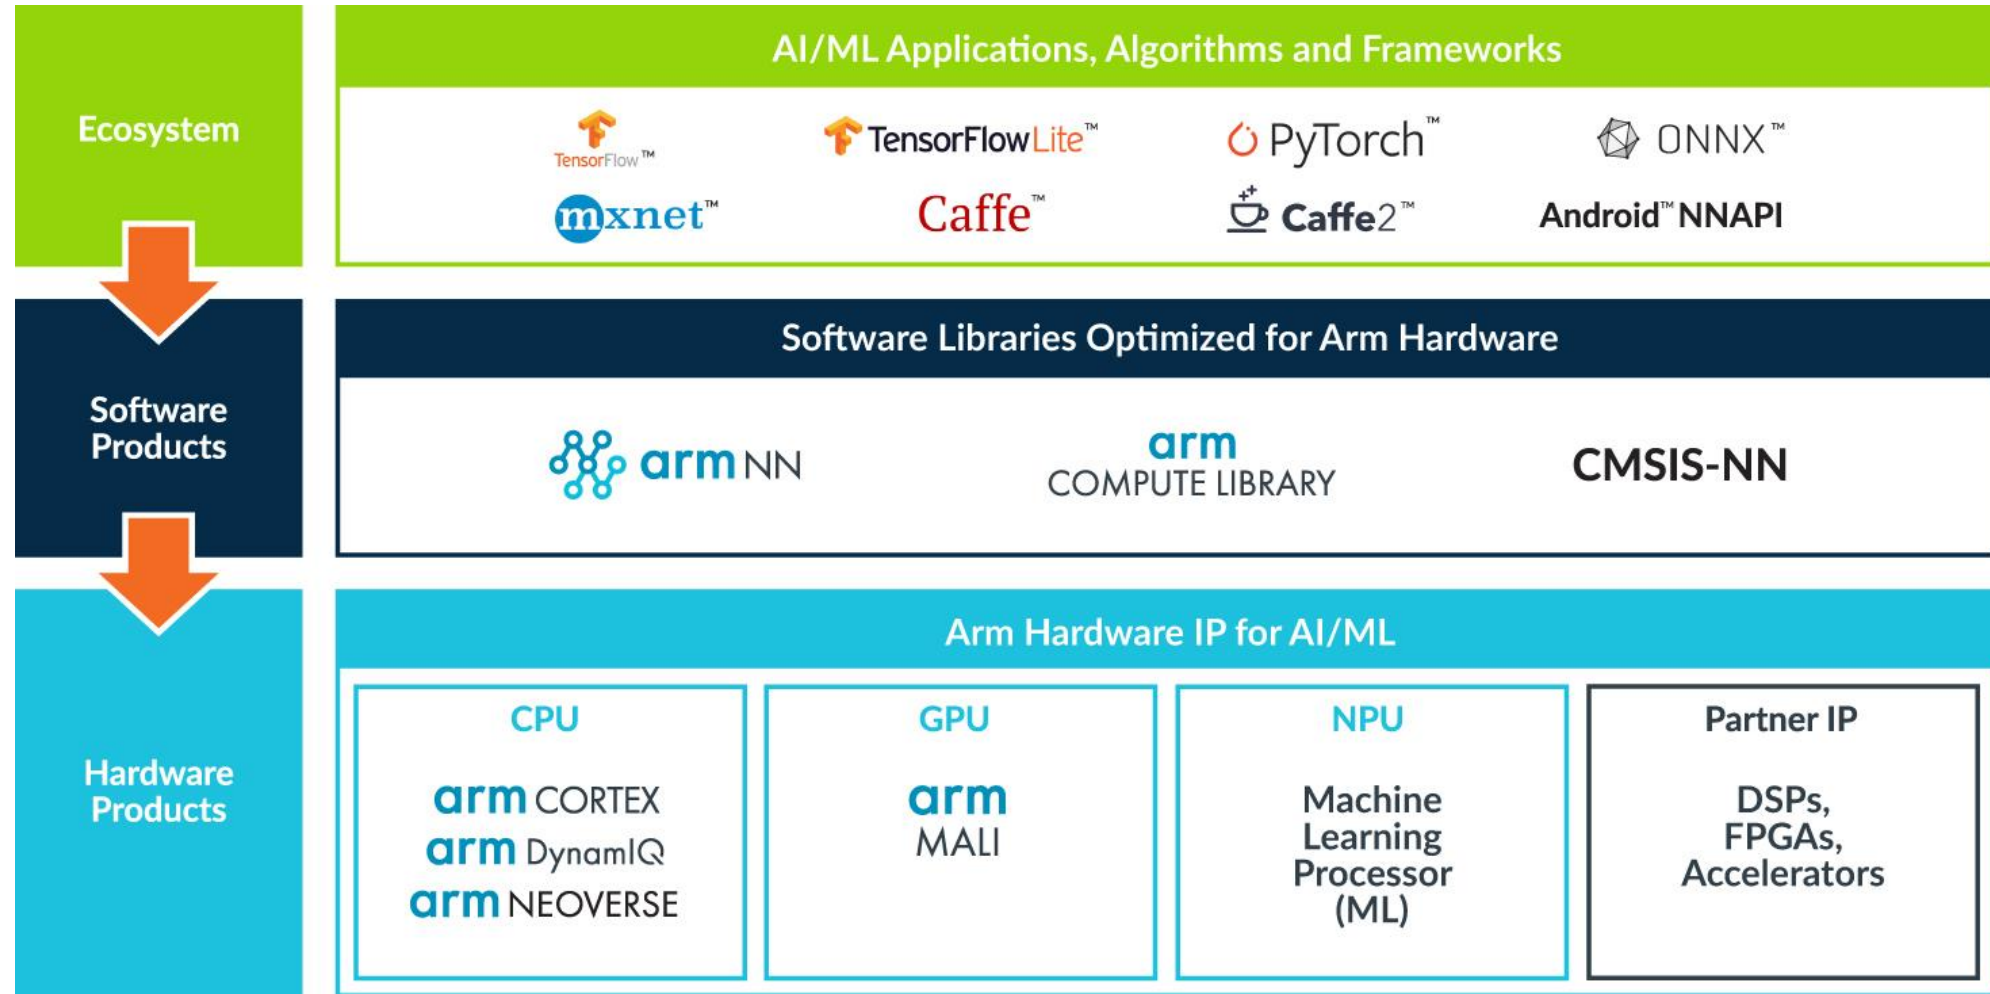
Compilers and Libraries

- Arm Compiler for Linux
 - LLVM based compiler optimized for Arm
 - C/C++ & Fortran
- GCC
 - Optimized for Arm
- Maths libraries
 - LAPACK, BLAS, FFT, SpMV, SpMM
 - Transcendentals
 - Micro-architecture optimized

Parallel Tools

- DDT: Debugger
 - C/C++, Fortran and Python
 - MPI, OpenMP, CUDA, OpenACC
- MAP: Performance Profiler
 - Scalable sampling-based profiler
- ArmIE: Instruction Emulator
 - Emulate unsupported instructions
 - SVE

Machine Learning and Artificial Intelligence

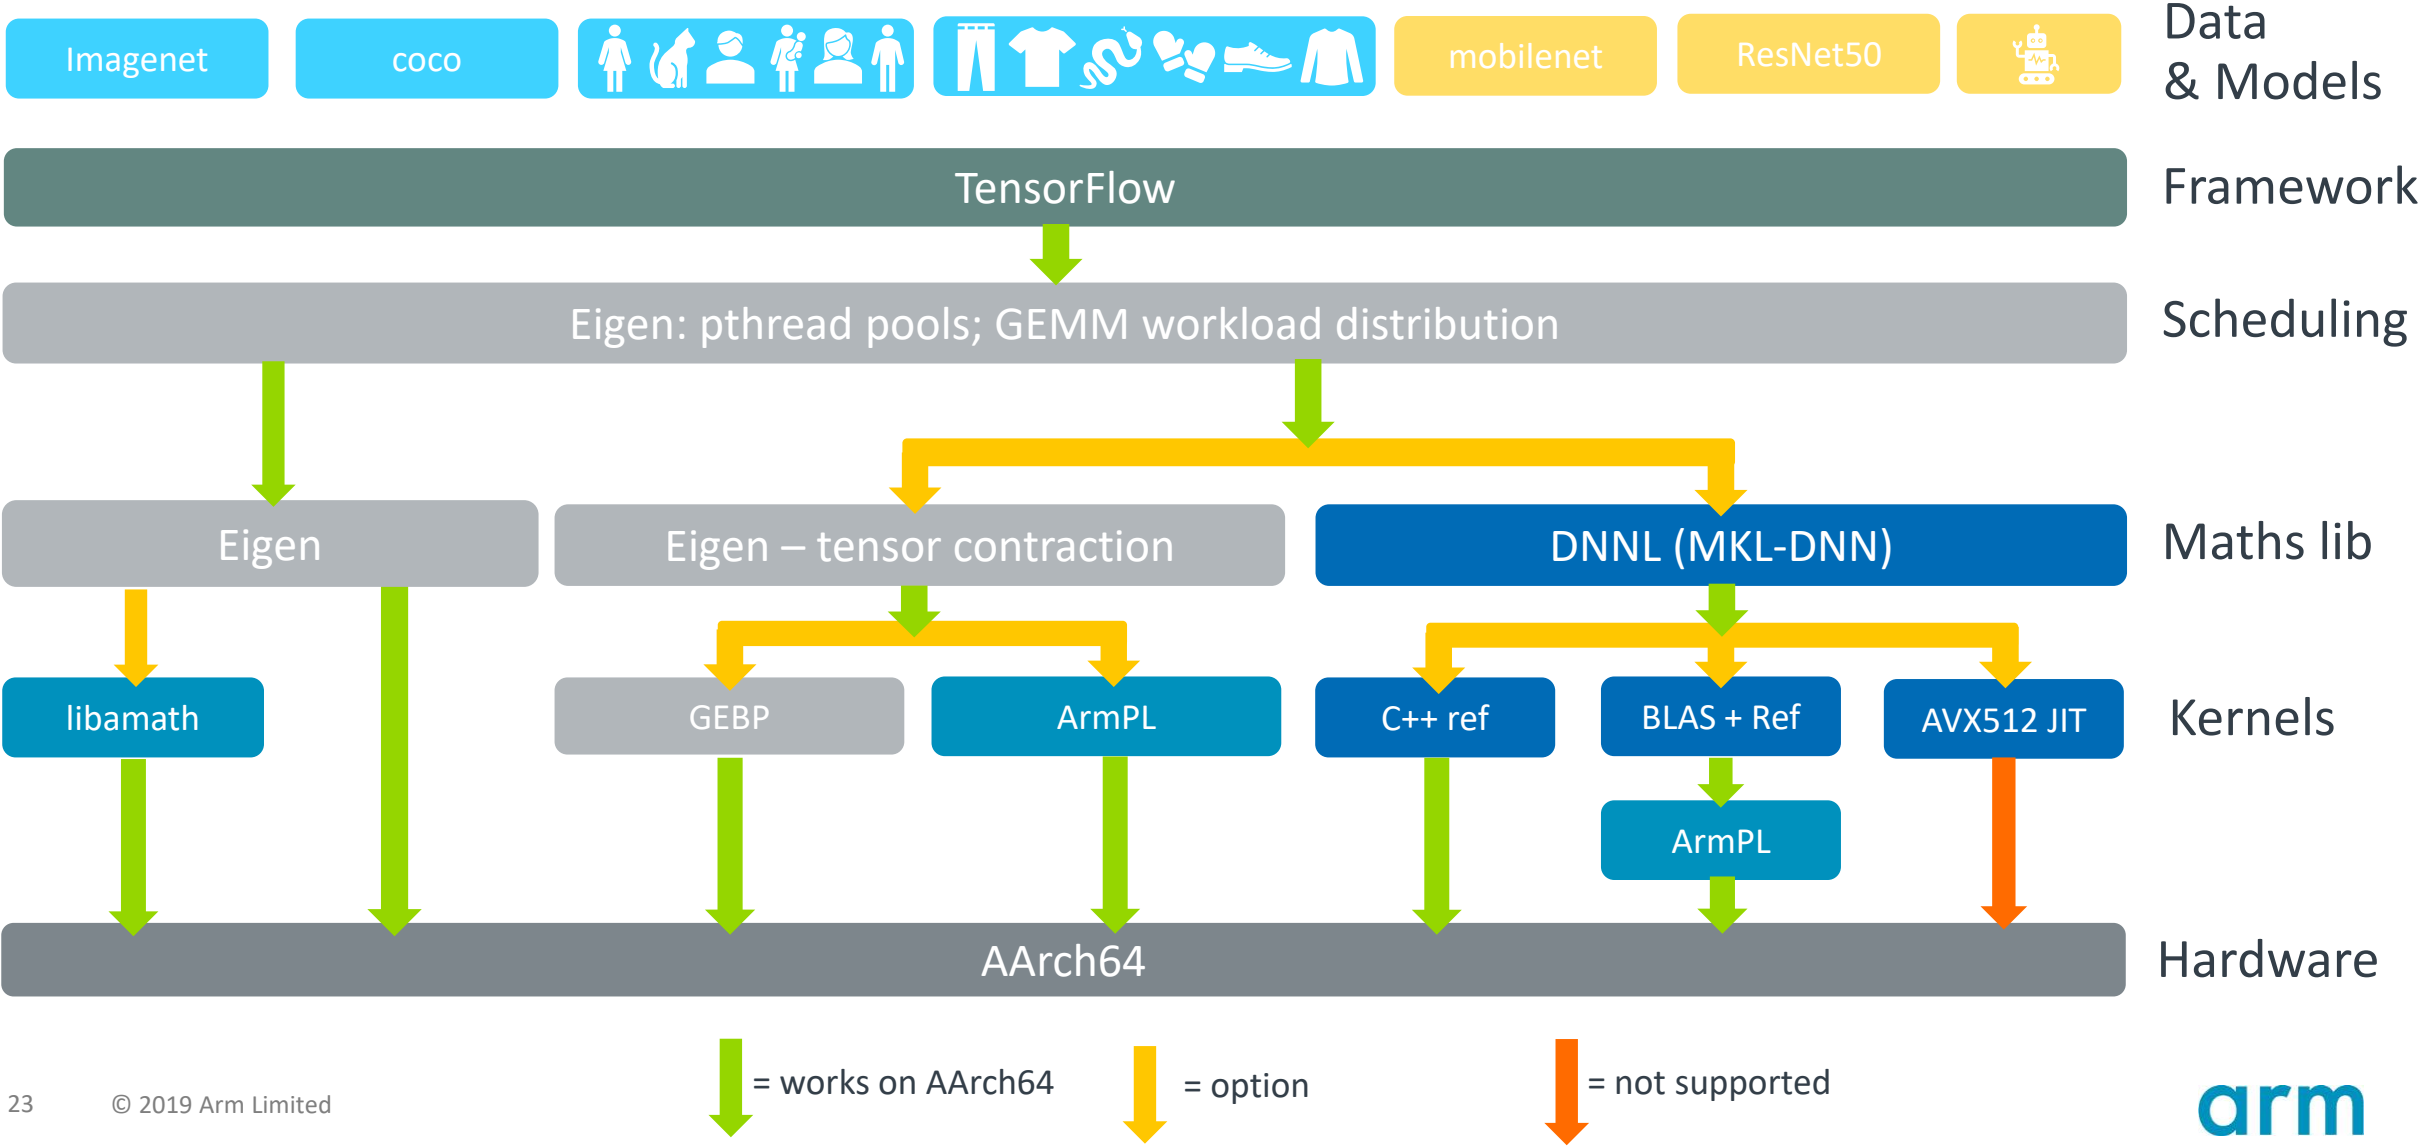


ML Frameworks on AArch64

On-CPU server-scale ML workloads

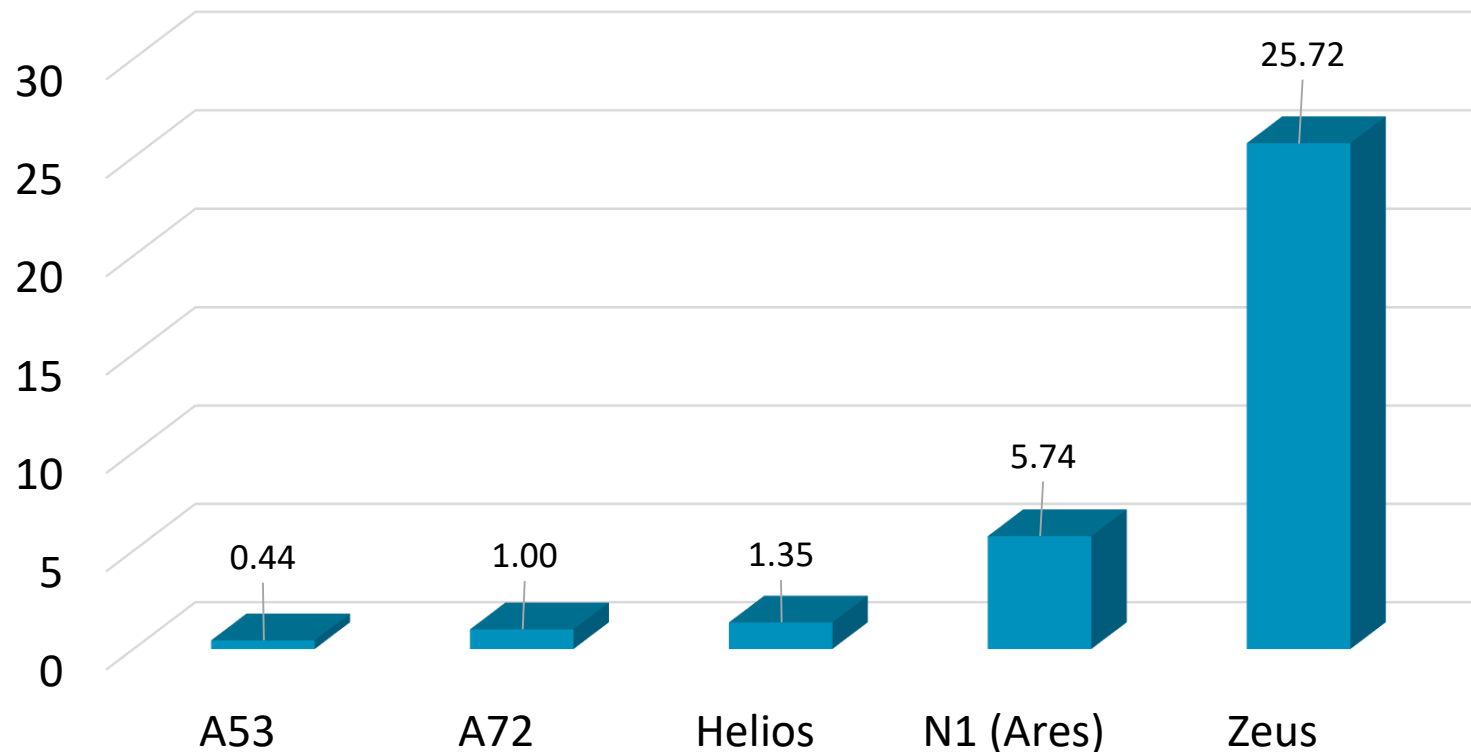
- Leading frameworks and dependencies built on AArch64
 - TensorFlow: <https://gitlab.com/arm-hpc/packages/-/wikis/packages/tensorflow>
 - PyTorch: <https://gitlab.com/arm-hpc/packages/-/wikis/packages/pytorch>
 - MXNET: <https://gitlab.com/arm-hpc/packages/-/wikis/packages/mxnet>
- OS Docker tools for TensorFlow part of *github.com/ARM-software/Tool-Solutions*
 - <https://github.com/ARM-software/Tool-Solutions/tree/master/docker/tensorflow-aarch64>
 - Compiler: GCC 9.2
 - Maths libraries: Arm Optimized Routines and OpenBLAS 0.3.6
 - Python3 environment built from CPython 3.7 and containing:
 - NumPy 1.17.1
 - TensorFlow 1.15
 - TensorFlow Benchmarks
- Focus has been on TensorFlow, and the maths libraries
 - Inference applications
 - Many-core systems
 - Significant GEMM, and vector maths, work

TensorFlow and maths libraries on AArch64



Increasing ML performance over CPU generations

Int8 GEMM kernel performance
(normalized to A72)



A72

2x ML performance
improvement over Cortex-A53

Helios

>3x ML performance
improvement over Cortex-A53
(First Multi-threaded CPU)

N1

>5x ML performance
improvement over Cortex-A72
(PPA leadership & ML
enhancements)

Zeus

>25x ML performance
improvement over Cortex-A72
(Breakthrough ML performance)


arm

NEON:
General Purpose Vector
Instructions



Arm NEON Vector Units

- SIMD Vector Extensions
 - Advanced Single Instruction Multiple Data
 - Fixed width at 128-bit
- As of Armv8-a
 - 31x 64-bit general-purpose registers
 - The 32-bit W register is lower half of 64-bit X register
 - 32x 128-bit floating-point registers
 - D is lower 64 bits of 128-bit Q registers
- Example:
 - fadd v0.4s, v0.4s, v1.4s
 - Addition of 4 x 32-bit floats
 - 128-bit NEON/32-bit Int = 4 Lanes

64-bit register layout



SPSR: Saved Program State Register



Arm NEON Vector Instructions

- Comprehensive set of vector operations
 - Loads, stores and maths operations
 - Scalar and floating point

$$\text{FADD} \quad \underbrace{\text{<Vd>}}_{\text{Destination}} \cdot \underbrace{\text{<T>}}_{\text{Type}}, \quad \underbrace{\text{<Vn>}}_{\text{Operand 1}} \cdot \underbrace{\text{<T>}}_{\text{Type}}, \quad \underbrace{\text{<Vm>}}_{\text{Operand 2}} \cdot \underbrace{\text{<T>}}_{\text{Type}}$$

- <Vd> - Destination register, <T> - Type e.g. 4S or 2D

$$\text{FADD} \quad \text{V0} \cdot 4\text{S}, \quad \text{V0} \cdot 4\text{S}, \quad \text{V1} \cdot 4\text{S}$$

- Add 4 single precision floating point values from V0 to V1 and store in V0
 - V0 += V1

Programming NEON

```
.LBB0_4:
    ldp     q0, q1, [x10, #-16]
    subs    x11, x11, #8
    fadd     v0.4s, v0.4s, v0.4s
    fadd     v1.4s, v1.4s, v1.4s
    stp     q0, q1, [x10, #-16]
    add     x10, x10, #32
    b.ne    .LBB0_4
// %bb.5:
    cmp     x9, x8
    b.eq    .LBB0_8
.LBB0_6:
    add     x10, x1, x9, lsl #2
    sub     x8, x8, x9
.LBB0_7:
    ldr     s0, [x10]
    subs    x8, x8, #1
    fadd     s0, s0, s0
    str     s0, [x10], #4
    b.ne    .LBB0_7
.LBB0_8:
    mov     w0, wzr
    ret
```

```
for (int i=0; i < n; ++i) {
    a[i] = 2.0 * a[i];
}
```

- .LBB0_4: Start with vector loop (and unroll)
 - Load 8 x 32-bit values (into 2 x 128-bit registers)
 - Subtract 8 from loop counter
 - 2x NEON add instructions (register to itself => 2.0*a[i])
 - Store pair of 128-bit registers back
 - Update array offsets
 - Loop if >=8 iterations left
- .LBB0_7: Remainder (fewer than 8 iterations left)
 - Load a single scalar
 - Add it to itself
 - Store
 - Loop if iterations left

Ease of Use for NEON

Where to start

NEON Intrinsics (ACLE)

- `#include arm_neon.h`
 - Header file of NEON intrinsics
- Map to assembly types and instruction names

```
float32x4_t va = vld1q_f32(&a[i]);  
va = vmulq_n_f32(va, 2.0);  
vst1q_f32(&a[i], va)
```

- Load 4x 32-bit floats into `va` from `a[i]`
- Multiply the floats in `va` by 2.0
- Store contents of `va` back into `a[i]`

Auto Vectorisation & Libraries

- Not everyone wants to hand code assembly
- Compilers will generate vector code
 - Generally at optimization levels `> -O2`
 - Supported in GCC, LLVM, Cray, Arm compiler
 - Vectorisation reports will inform on success
- Vectorised libraries
 - Such as ArmPL maths library

Limitations of NEON

- NEON is firstly only 128-bit
 - Not much use to HPC / Scientific Computing
- Want bigger vectors
 - To expand to 256-bit or 512-bit we would need separate instructions
 - Arm like to offer flexibility to customers - Different vector lengths
 - However it is a RISC architecture (32-bit instruction encoding)
- Suffers from same drawbacks as other vector implementations (AVX)
 - Difficulty to autovectorise
 - Remainder loops

arm

SVE:
Today's new
Vectorisation Paradigm

SVE: Scalable Vector Extension

- **SVE is Vector Length Agnostic (VLA)**
 - Vector Length (VL) is a hardware implementation choice from 128 up to 2048 bits.
 - New programming model allows software to scale dynamically to available vector length.
 - No need to define a new ISA, rewrite or recompile for new vector lengths.
- **SVE is not an extension of Advanced SIMD (*aka* Neon)**
 - A separate, optional extension with a new set of instruction encodings.
 - Initial focus is HPC and general-purpose server, not media/image processing.
- **SVE begins to tackle traditional barriers to auto-vectorization**
 - Software-managed speculative vectorization allows uncounted loops to be vectorized.
 - In-vector serialised inner loop permits outer loop vectorization in spite of dependencies.

How can you program when the vector length is unknown?

SVE provides features to enable VLA programming from the assembly level and up

	1	2	3	4
	5	5	5	5
+				
<i>pred</i>	1	0	1	0
=	6	2	8	4

Per-lane predication

Operations work on individual lanes under control of a predicate register.

for	(i = 0; i < n; ++i)
<i>INDEX i</i>	n-2 n-1 n n+1
<i>CMPLT n</i>	1 1 0 0

Predicate-driven loop control and management

Eliminate scalar loop heads and tails by processing partial vectors.



	1	2		
	1	2	0	0
+				
<i>pred</i>	1	1	0	0

Vector partitioning & software-managed speculation

First Faulting Load instructions allow memory accesses to cross into invalid pages.

SVE Registers


- **Scalable vector registers**
 - Z0-Z31 extending NEON's 128-bit v0-v31.
 - Packed DP, SP & HP floating-point elements.
 - Packed 64, 32, 16 & 8-bit integer elements.
- **Scalable predicate registers**
 - P0-P15 predicates for loop / arithmetic control.
 - $1/8^{\text{th}}$ size of SVE registers (1 bit / byte).
 - FFR first fault register for software speculation.



Predicates: Active Lanes vs Inactive Lanes

Predicate registers track lane activity

- 16 predicate registers (P0-P15)
- 1 predicate bit per 8 vector bits (lowest predicate bit per lane is significant)
- On **load**, active elements update the destination
- On **store**, inactive lanes leave destination unchanged (p0/m) or set to 0's (p0/z)



SVE Predicate condition flags

SVE is a *predicate-centric* architecture

- Predicates are central, not an afterthought
- Support complex nested conditions and loops.
- Predicate generation also sets condition flags.
- Reduces vector loop management overhead.

Overloading the A64 NZCV condition flags

Flag	SVE	Condition
N	First	Set if first active element is true
Z	None	Set if no active element is true
C	!Last	Set if last active element is false
V		Scalarized loop state, else zero


Reuses the A64 conditional instructions

- Conditional branches B.EQ → B.NONE
- Conditional select, set, increment, etc.

Condition Test	A64 Name	SVE Alias	SVE Interpretation
Z=1	EQ	NONE	No active elements are true
Z=0	NE	ANY	Any active element is true
C=1	CS	NLAST	Last active element is not true
C=0	CC	LAST	Last active element is true
N=1	MI	FIRST	First active element is true
N=0	PL	NFRST	First active element is not true
C=1 & Z=0	HI	PMORE	More partitions: some active elements are true but not the last one
C=0 Z=1	LS	PLAST	Last partition: last active element is true or none are true
N=V	GE	TCONT	Continue scalar loop
N!=V	LT	TSTOP	Stop scalar loop

SVE supports vectorization in complex code

Right from the start, SVE was engineered to handle codes that usually won't vectorize



Gather-load and scatter-store

Loads a single register from several non-contiguous memory locations.

$$\begin{array}{ccccccc} \boxed{1} & + & \boxed{2} & + & \boxed{3} & + & \boxed{4} & = \\ \boxed{1} & + & \boxed{2} & & \boxed{3} & + & \boxed{4} & \\ & & = & & = & & & \\ \boxed{3} & & + & & \boxed{7} & & = \end{array}$$

Extended floating-point horizontal reductions

In-order and tree-based reductions trade-off performance and repeatability.

arm

Questions and
Break



arm




Part 2: Going Forward with SVE

Quick Recap

- SVE enables Vector Length Agnostic (VLA) programming
- VLA enables portability, scalability, and optimization
- Predicates control which operations affect which vector lanes
 - Predicates are not bitmasks
 - You can think of them as dynamically resizing the vector registers
- The actual vector length is set by the CPU architect
 - Any multiple of 128 bits up to 2048 bits
 - May be dynamically reduced by the OS or hypervisor
- SVE was designed for HPC and can vectorize complex structures
- Many open source and commercial tools currently support SVE

```
for (i = 0; i < n; ++i)
    INDEX i
    CMPLT n
```

	n-2	n-1	n	n+1
	1	1	0	0



arm

VLA Programming

Vector Length Agnostic

Vector Length Agnostic Programming

A paradigm shift for developers

Advantages

- Not thinking about vector length
 - Rather just vectorization
- No peel/remainder loops
 - All handled by predication
- Key is loop structures
 - Predicates are powerful

Considerations

- Should not be writing fixed width
 - Applies to data structures and instructions
 - More portable for different hardware
- Can the compiler identify loop structure?
 - Generate predicated instructions
- However: VLA **may** be slower
 - Cost of generating predicates
 - Near empty loops

How do you count by vector width?

No need for multi-versioning: one increment to rule all vector sizes

```
ld1w z1.s, p0/z, [x0,x4,ls1 2]      // p0:z1=x[i]
ld1w z2.s, p0/z, [x1,x4,ls1 2]      // p0:z2=y[i]
fm1a z2.s, p0/m, z1.s, z0.s         // p0?z2+=x[i]*a
st1w z2.s, p0, [x1,x4,ls1 2]        // p0?y[i]=z2

incw x4                             // i+=(VL/32)
```

“Increment `x4` by the number of 32-bit lanes (**w**) that fit in a VL.”

Initialization when vector length is unknown

- Vectors cannot be initialized from compile-time constant, so...
 - INDEX `Zd.S, #1, #4` : `Zd = [1, 5, 9, 13, 17, 21, 25, 29]`
- Predicates cannot be initialized from memory, so...
 - PTRUE `Pd.S, MUL3` : `Pd = [T, T, T, T, T, T, F, F]`
- Vector loop increment and trip count are unknown at compile-time, so...
 - INCD `Xi` : increment scalar `Xi` by # of 64b dwords in vector
 - WHILELT `Pd.D, Xi, Xe` : next iteration predicate `Pd = [while i++ < e]`
- Vectors stores to stack must be dynamically allocated and indexed, so...
 - ADDVL `SP, SP, #-4` : decrement stack pointer by (`4*VL`)
 - STR `Zi, [SP, #3, MUL VL]` : store vector `Z1` to address (`SP+3*VL`)

Vectorizing A Scalar Loop With ACLE

$a[:] = 2.0 * a[:]$

Original Code

```
for (int i=0; i < N; ++i) {  
  
    a[i] = 2.0 * a[i];  
  
}
```

128-bit NEON vectorization

```
int i;  
  
// vector loop  
for (i=0; (i<N-3) && (N&~3); i+=4) {  
    float32x4_t va = vld1q_f32(&a[i]);  
    va = vmulq_n_f32(va, 2.0);  
    vst1q_f32(&a[i], va)  
}  
  
// drain loop  
for (; i < N; ++i)  
    a[i] = 2.0 * a[i];
```

This is NEON,
not SVE!

Vectorizing A Scalar Loop With ACLE

`a[:] = 2.0 * a[:]`

```
for (int i=0; i < N; ++i) {  
    a[i] = 2.0 * a[i];  
}
```

128-bit NEON vectorization

```
int i;  
  
// vector loop  
for (i=0; (i<N-3) && (N&~3); i+=4) {  
    float32x4_t va = vld1q_f32(&a[i]);  
    va = vmulq_n_f32(va, 2.0);  
    vst1q_f32(&a[i], va)  
}  
  
// drain loop  
for (; i < N; ++i)  
    a[i] = 2.0 * a[i];
```

SVE vectorization

```
for (int i = 0 ; i < N; i += svcntw() )  
{  
    svbool_t Pg = svwhilelt_b32(i, N);  
    svfloat32_t va = svld1(Pg, &a[i]);  
    va = svmul_x(Pg, va, 2.0);  
    svst1(Pg, &a[i], va);  
}
```

Vectorizing A Scalar Loop With ACLE

$a[:] = 2.0 * a[:]$

```
for (int i=0; i < N; ++i) {  
    a[i] = 2.0 * a[i];  
}
```

SVE vectorization

```
for (int i = 0 ; i < N; i += svcntw() )  
{  
    svbool_t Pg = svwhilelt_b32(i, N);  
    svfloat32_t va = svld1(Pg, &a[i]);  
    va = svmul_x(Pg, va, 2.0);  
    svst1(Pg, &a[i], va);  
}
```

Assembly

armclang -march=armv8.2-a+sve -O3 -S sve2.c

```
cmp    w0, #1  
b.lt   .LBB0_3  
  
mov    w8, wzr  
.LBB0_2:  
whilelt p0.s, w8, w0  
sxtw   x9, w8  
ld1w   { z0.s }, p0/z, [x1, x9, lsl #2]  
incw   x8  
cmp    w8, w0  
fmul   z0.s, p0/m, z0.s, #2.0  
st1w   { z0.s }, p0, [x1, x9, lsl #2]  
b.lt   .LBB0_2  
.LBB0_3:  
ret
```

SVE Gives You More

- SVE is really powerful (mainly due to predicates)
 - Compilers can exploit this power
 - Autovectorisation getting much better
- Power is also being able to vectorise new things
 - Previously hard to vectorise
 - Mapping IF statements to predicates
- Shown in the following real-world example
 - Users code had a 'mask' in the inner loop
 - Previously prevented vectorisation

More Complex Example

Code

```
void foo(char* mask,
         double * result, int n)
{
    for (int x = 0; x < n; x++)
        if(mask[x] == 0)
            result[x] += 2.0;
}
```

```
armclang -S -march=armv8.2-a+sve -O3 sve.c
```

Assembly

```
.LBB0_7:
    mov     x9, xzr
    whilelo p0.d, xzr, x8
    mov     z0.d, #0
    fmov     z1.d, #2.000000000

.LBB0_8:
    ld1b     { z2.d }, p0/z, [x0, x9]
    cmpeq     p0.d, p0/z, z2.d, z0.d
    ld1d     { z2.d }, p0/z, [x1, x9, lsl #3]
    fadd     z2.d, z2.d, z1.d
    st1d     { z2.d }, p0, [x1, x9, lsl #3]
    incd     x9
    whilelo  p0.d, x9, x8
    b.mi     .LBB0_8

.LBB0_9:
    ret
```

Annotations

Generate initial loop predicate
Generate initial register values

Load mask with loop predicate
Generate predicate mask[i]==0
Load result with mask predicate
Addition with mask predicate
Store with mask predicate
Loop counter increment
Generate new loop predicate
Loop

arm

Toolchain Support

Getting Help

How to Make Use of SVE

Using the tools that are available

- Many ways to use SVE in your application
 - Different levels of effort, reward and portability
- SVE libraries
 - Such as SVE ArmPL for SVE maths operations, just link the correct version
- Auto-vectorisation
 - Compiler spots vectorisation and exploits it, compiler reports to inform
 - `-march=armv8-a+sve`
- Pragmas
 - Help and guide auto-vectorisation (e.g. OMP SIMD, IVDEP)
- Intrinsics
 - ACLE for SVE - friendly interface into SVE instructions
- Assembly
 - Hand code your SVE instructions

Arm Compiler for Linux user-space

Commercial C/C++/Fortran compiler with best-in-class performance



Compilers tuned for Scientific Computing and HPC



Latest features and performance optimizations



Commercially supported by Arm

Tuned for Scientific Computing, HPC and Enterprise workloads

- Processor-specific optimizations for various server-class Arm-based platforms
- Optimal shared-memory parallelism using latest Arm-optimized OpenMP runtime

Linux user-space compiler with latest features

- C++ 14 and Fortran 2003 language support with OpenMP 4.5
- Support for Armv8-A and SVE architecture extension
- Based on LLVM and Flang, leading open-source compiler projects

Commercially supported by Arm

- Available for a wide range of Arm-based platforms running leading Linux distributions – RedHat, SUSE and Ubuntu

Arm Compiler – Auto-vectorization

LLVM9

- Arm pulls all relevant cost models and optimizations into the downstream codebase.
- Auto-vectorization via LLVM **vectorizers**:
 - Use cost models to drive decisions about what code blocks can and/or should be vectorized.
 - Since October 2018, two different vectorizers used from LLVM: [Loop Vectorizer](#) and [SLP Vectorizer](#).
- Loop Vectorizer support for SVE and NEON:
 - Loops with unknown trip count
 - Runtime checks of pointers
 - Reductions
 - Inductions
 - “If” conversion
 - Pointer induction variables
 - Reverse iterators
 - Scatter / gather
 - Vectorization of mixed types
 - Global structures alias analysis

New: ACFL Vectorisation Reports

```
`-fsave-optimization-record` & `arm-opt-report file.opt.yaml`
```

Vectorized
4x 32-bit lanes
1-way
interleaving

Fully
unrolled

All 3 instances of `foo()` were inlined

```
$ armclang -O3 or.c -c -o or.o -fsave-optimization-record -march=armv8-a+sve
$ arm-opt-report or.opt.yaml
< or.c
1      | void bar();
2      | void foo() { bar(); }
3      |
4      | void Test(int *res, int *c, int *d, int *p, int n) {
5      |     int i;
6      |
7      |     #pragma clang loop vectorize(assume_safety)
8      |     for (i = 0; i < 1600; i++) {
9      |         res[i] = (p[i] == 0) ? res[i] : res[i] + d[i];
10     |     }
11     |
12     |     for (i = 0; i < 16; i++) {
13     |         res[i] = (p[i] == 0) ? res[i] : res[i] + d[i];
14     |     }
15     |
16     |     foo();
17     |
18     |     foo(); bar(); foo();
19     |     ^
20     |     ^
21     | }
```

Arm Performance Libraries

Optimized BLAS, LAPACK and FFT



Commercially supported
by Arm



Best in class performance



Validated with
NAG test suite

Commercial 64-bit Armv8-A math libraries

- Commonly used low-level math routines - BLAS, LAPACK and FFT
- Provides FFTW compatible interface for FFT routines
- Sparse linear algebra and batched BLAS support
- libamath gives high-performing math.h functions implementations

Best-in-class serial and parallel performance

- Generic Armv8-A optimizations by Arm
- Tuning for specific platforms like Marvell ThunderX2 in collaboration with silicon vendors

Validated and supported by Arm

- Available for a wide range of server-class Arm-based platforms
- Validated with NAG's test suite, a de-facto standard

SVE Compiler Support



Compiler	Assembly / Disassembly	Inline Assembly	ACLE	Auto- vectorization	Math Libraries
Arm Compiler for HPC	SVE + SVE2	SVE + SVE2	SVE + SVE2	SVE+ SVE2	SVE
LLVM/Clang	SVE + SVE2	SVE + SVE2	SVE + SVE2 in LLVM 10	SVE + SVE2 in LLVM 11	
GNU	SVE + SVE2	SVE + SVE2	SVE + SVE2 in GNU 10	SVE now SVE2 in GNU10	

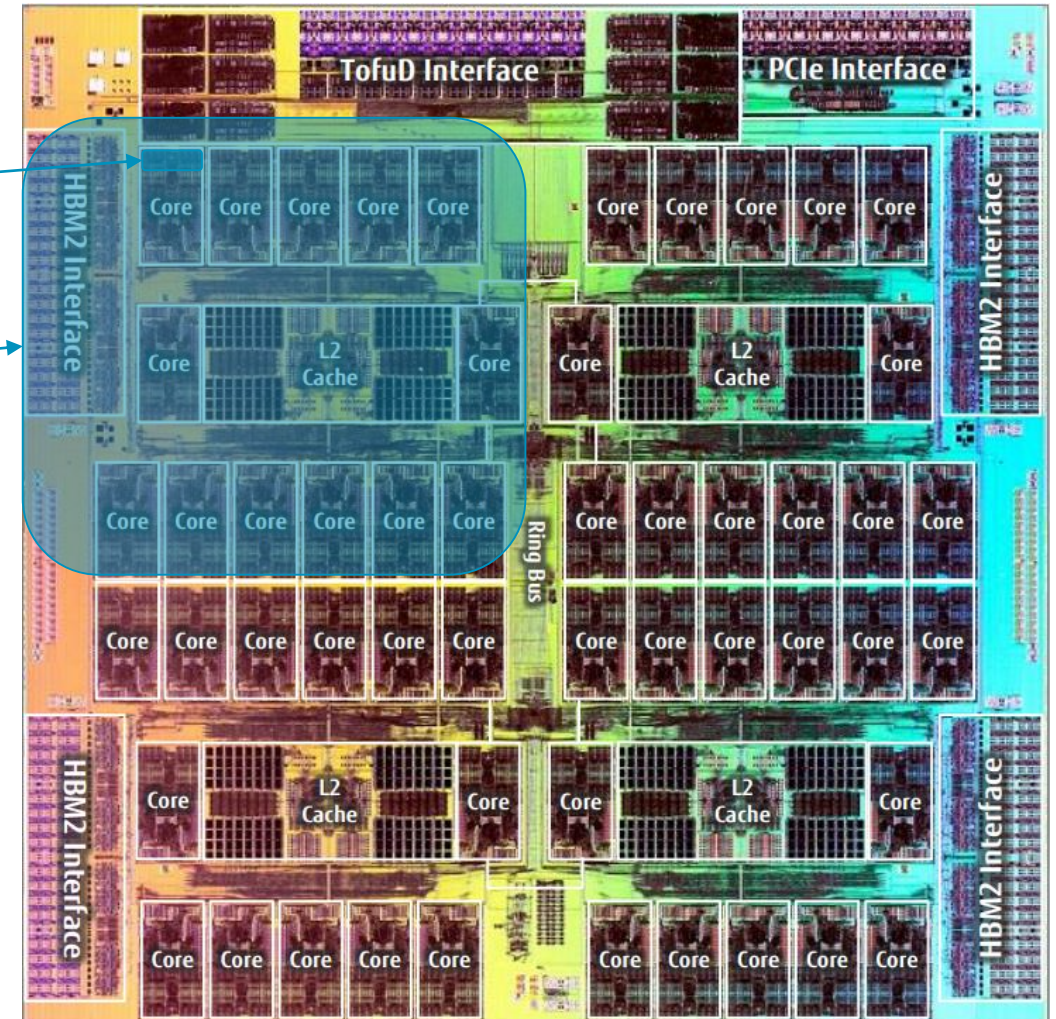
arm

Realising SVE

The First SVE CPU - Available ~Now

Fujitsu A64FX

- Custom Arm design by Fujitsu
- 512-bit SVE
- 4 Core Memory Groups
 - 12 cores + 1 OS core
 - 8 GB High Bandwidth Memory (HBM)
- 1 socket / node, 2 nodes / blade
 - Direct water cooling



A64FX Now in Top500 - #159

System	Year	Vendor	Cores	Rmax (GFlop/s)	Rpeak (GFlop/s)
A64FX prototype - Fujitsu A64FX, Fujitsu A64FX 48C 2GHz, Tofu interconnect D	2019		36,864	1,999,500	2,359,296

Green500 - #1

TOP500		System	Cores	Rmax (TFlop/s)	Power (kW)	Power Efficiency (GFlops/watts)
Rank	Rank					
1	159	A64FX prototype - Fujitsu A64FX, Fujitsu A64FX 48C 2GHz, Tofu interconnect D , Fujitsu Numazu Plant Japan	36,864	1,999.5	118	16.876



arm

Post SVE: The Vector Instructions of the Next Generation

FMMLA: High Performance Matrix Multiplication

- Added to Armv8.6
 - NEON and SVE instructions
 - FMMLA instructions for FP (SVE)

FMMLA <Zda>.S, <Zn>.S, <Zm>.S

FMMLA <Zda>.D, <Zn>.D, <Zm>.D

- 2x2 matrix multiplication
 - Works on multiple of 'vector granules'
 - 2x2xFP32 = 128-bit granules
 - Assumes vector length is multiple
- May require layout transformations
 - Outer loop to minimise cost
- Accelerated libraries

0	1
2	3

 +=


0	1
2	3

 X

0	1
2	3


Dest (D) Left (L) Right (R)
2x2xFP32 2x2xFP32 2x2xFP32

$$\begin{aligned} D[0] &+= (L[0] * R[0]) + (L[1] * R[1]) \\ D[1] &+= (L[0] * R[2]) + (L[1] * R[3]) \\ D[2] &+= (L[2] * R[0]) + (L[3] * R[1]) \\ D[3] &+= (L[2] * R[2]) + (L[3] * R[3]) \end{aligned}$$



New Data Type Support: BFloat16

- New addition to Armv8-A
 - Adds support for BF16
- Instructions for NEON and SVE
 - Including:
 - **BFDOT**: Dot Product (1x2)x(2x1)
 - **BFMMLA**: Mat Multiply (2x4)x(4x2)
- Significant performance gains
 - ML training and inference workloads
- Supported in Arm libraries
 - Arm NN and Arm Compute Libraries



Scalable Vector Extensions V2 (SVE2)

SVE for non HPC markets



Built on SVE



Improved scalability



Vectorization of
more workloads

- Built on the SVE foundation.
 - Scalable vectors with hardware choice from 128 to 2048 bits.
 - Vector-length agnostic programming for “write once, run anywhere”.
 - Tackles some obstacles to compiler auto-vectorisation.
- Scaling single-thread performance to exploit long vectors.
 - SVE2 adds NEON™-style fixed-point DSP/multimedia plus other new features.
 - Performance parity and beyond with classic NEON DSP/media SIMD.
 - Tackles further obstacles to compiler auto-vectorization.
- Enables vectorization of a wider range of applications than SVE.
 - Multiple use cases in Client, Edge, Server and HPC.
 - DSP, Codecs/filters, Computer vision, Photography, Game physics, AR/VR,
 - Networking, Baseband, Database, Cryptography, Genomics, Web serving.
 - Improves competitiveness of Arm-based CPU vs proprietary solutions.
 - Reduces s/w development time and effort.

SVE2 Instructions Add:

What's new


- Thorough support for fixed-point DSP arithmetic
 - (traditional Neon DSP/Media processing, complex numbers arithmetic for LTE)
- Multi-precision arithmetic
 - (bignum, crypto)
- Non-temporal gather/scatter
 - (HPC, sort)
- Enhanced permute and bitwise permute instructions
 - (CV, FIR, FFT, LTE, ML, genomics, cryptanalysis)
- Histogram acceleration support
 - (CV, HPC, sort)
- String processing acceleration support
 - (parsers)
- (optional) Cryptography support instructions for AES, SM4, SHA standards
 - (encryption)

Example: Widening and Narrowing


NEON vs SVE2

NEON

UADDL Vd.2D, Vn.2S, Vm.2S




UADDL2 Vd.2D, Vn.4S, Vm.4S




- NEON uses high/low half of vector
- Expensive for large vector lengths
 - >> 128-bit
- SVE2 uses odd/even half of vector
- Bottom and top
- Happens 'in-lane'

SVE2

UADDLB Zd.D, Zn.S, Zm.S



UADDLT Zd.D, Zn.S, Zm.S



Transactional Memory Extension (TME)

Scalable Thread-Level Parallelism (TLP) for multi-threaded applications



Hardware Transactional
Memory



Improved scalability



Simpler software design

- Hardware Transactional Memory (HTM) for the Arm architecture.
 - Improved competitiveness with other architectures that support HTM.
 - Strong isolation between threads.
 - Failure atomicity.
- Scaling multi-thread performance to exploit many-core designs.
 - Database.
 - Network dataplane.
 - Dynamic web serving.
- Simplifies software design for massively multi-threaded code.
 - Supports Transactional Lock Elision (TLE) for existing locking code.
 - Low-level concurrent access to shared data is easier to write and debug.

arm

Conclusion

Vector Instruction Sets on Arm

- Arm is new to HPC but has a compelling justification and heritage
- Arm has a constantly evolving set of vector instructions
 - From NEON to SVE and beyond
- Designed to give flexibility to hardware designers
 - Maximise ability to differentiate
 - Still present a consistent end-user experience
- Designed for real use-cases for improved performance
 - Started with HPC workloads
 - Moving to AI / ML server

The ARM logo is displayed in a white, lowercase, sans-serif font. The letters are bold and modern, with the 'a' and 'm' having a slightly rounded, friendly appearance. The logo is centered horizontally in the upper left portion of the slide.

The Arm trademarks featured in this presentation are registered trademarks or trademarks of Arm Limited (or its subsidiaries) in the US and/or elsewhere. All rights reserved. All other marks featured may be trademarks of their respective owners.

www.arm.com/company/policies/trademarks