




Arm® Cortex®-A76 Software Optimization Guide

Version 8.0



Arm® Cortex-A76

Software Optimization Guide

Copyright © 2019 Arm Limited (or its affiliates). All rights reserved.

Release Information

The following changes have been made to this Software Optimization Guide.

Non-Confidential Proprietary Notice

This document is protected by copyright and other related rights and the practice or implementation of the information contained in this document may be protected by one or more patents or pending patent applications. No part of this document may be reproduced in any form by any means without the express prior written permission of Arm. **No license, express or implied, by estoppel or otherwise to any intellectual property rights is granted by this document unless specifically stated.**

Your access to the information in this document is conditional upon your acceptance that you will not use or permit others to use the information for the purposes of determining whether implementations infringe any third party patents.

THIS DOCUMENT IS PROVIDED “AS IS”. ARM PROVIDES NO REPRESENTATIONS AND NO WARRANTIES, EXPRESS, IMPLIED OR STATUTORY, INCLUDING, WITHOUT LIMITATION, THE IMPLIED WARRANTIES OF MERCHANTABILITY, SATISFACTORY QUALITY, NON-INFRINGEMENT OR FITNESS FOR A PARTICULAR PURPOSE WITH RESPECT TO THE DOCUMENT. For the avoidance of doubt, Arm makes no representation with respect to, and has undertaken no analysis to identify or understand the scope and content of, patents, copyrights, trade secrets, or other rights.

This document may include technical inaccuracies or typographical errors.

TO THE EXTENT NOT PROHIBITED BY LAW, IN NO EVENT WILL ARM BE LIABLE FOR ANY DAMAGES, INCLUDING WITHOUT LIMITATION ANY DIRECT, INDIRECT, SPECIAL, INCIDENTAL, PUNITIVE, OR CONSEQUENTIAL DAMAGES, HOWEVER CAUSED AND REGARDLESS OF THE THEORY OF LIABILITY, ARISING OUT OF ANY USE OF THIS DOCUMENT, EVEN IF ARM HAS BEEN ADVISED OF THE POSSIBILITY OF SUCH DAMAGES.

This document consists solely of commercial items. You shall be responsible for ensuring that any use, duplication or disclosure of this document complies fully with any relevant export laws and regulations to assure that this document or any portion thereof is not exported, directly or indirectly, in violation of such export laws. Use of the word “partner” in reference to Arm’s customers is not intended to create or refer to any partnership relationship with any other company. Arm may make changes to this document at any time and without notice.

If any of the provisions contained in these terms conflict with any of the provisions of any click through or signed written agreement covering this document with Arm, then the click through or signed written agreement prevails over and supersedes the conflicting provisions of these terms. This document may be translated into other languages for convenience, and you agree that if there is any conflict between the English version of this document and any translation, the terms of the English version of the Agreement shall prevail.

The Arm corporate logo and words marked with ® or ™ are registered trademarks or trademarks of Arm Limited (or its subsidiaries) in the US and/or elsewhere. All rights reserved. Other brands and names mentioned in this document may be the trademarks of their respective owners. Please follow Arm’s trademark usage guidelines at <http://www.arm.com/company/policies/trademarks>.

Copyright © 2019 Arm Limited or its affiliates. All rights reserved.

Arm Limited. Company 02557590 registered in England.

110 Fulbourn Road, Cambridge, England CB1 9NJ.

LES-PRE-20349

Confidentiality Status

This document is Non-Confidential. The right to use, copy and disclose this document may be subject to license restrictions in accordance with the terms of the agreement entered into by Arm and the party that Arm delivered this document to.

Unrestricted Access is an Arm internal classification.

Product Status

The information in this document is Final, that is for a developed product.

Web Address

<http://www.arm.com>.

Contents

1 About this document	6
1.1. References	6
1.2. Terms and Abbreviations	6
1.3. Scope	6
2 Introduction	7
2.1. Pipeline Overview	7
3 Instruction characteristics.....	8
3.1. Instruction tables	8
3.2. Legend for reading the utilized pipelines.....	8
3.3. Branch Instructions.....	8
3.4. Arithmetic and Logical Instructions	9
3.5. Move and Shift Instructions.....	10
3.6. Divide and Multiply Instructions.....	10
3.7. Saturating and Parallel Arithmetic Instructions	12
3.8. Miscellaneous Data-Processing Instructions	12
3.9. Load Instructions	13
3.10. Store Instructions	16
3.11. FP Data Processing Instructions.....	18
3.12. FP Miscellaneous Instructions	20
3.13. FP Load Instructions.....	21
3.14. FP Store Instructions.....	22
3.15. ASIMD Integer Instructions.....	24
3.16. ASIMD Floating-Point Instructions.....	28
3.17. ASIMD Miscellaneous Instructions	30
3.18. ASIMD Load Instructions	33
3.19. ASIMD Store Instructions.....	36
3.20. Cryptography Extensions	39
3.21. CRC	40
4 Special considerations	41
4.1. Dispatch Constraints.....	41
4.2. Dispatch Stall	41
4.3. Optimizing General-Purpose Register Spills and Fills.....	41
4.4. Optimizing Memory Copy	41
4.5. Load/Store Alignment.....	42

4.6. AES Encryption/Decryption 42

4.7. Region Based Fast Forwarding..... 43

4.8. Branch instruction alignment 43

4.9. FPCR self-synchronization..... 43

4.10. Special Register Access 44

4.11. Register Forwarding Hazards 45

4.12. IT Blocks..... 46

1 About this document

This document contains a guide to the Cortex-A55 micro-architecture with a view to aiding software optimization.

1.1. References

Reference	Document	Author	Title
1	DDI 0487	Arm	<i>Arm® Architecture Reference Manual, Armv8, for Armv8-A architecture profile</i>
2	100798	Arm	<i>Arm® Cortex®-A76 Core Technical Reference Manual</i>

1.2. Terms and Abbreviations

This document uses the following terms and abbreviations.

Term	Meaning
ALU	Arithmetic/Logical Unit
ASIMD	Advanced SIMD
Mop	Macro-Operation
Uop	Micro-Operation
VFP	Vector Floating Point

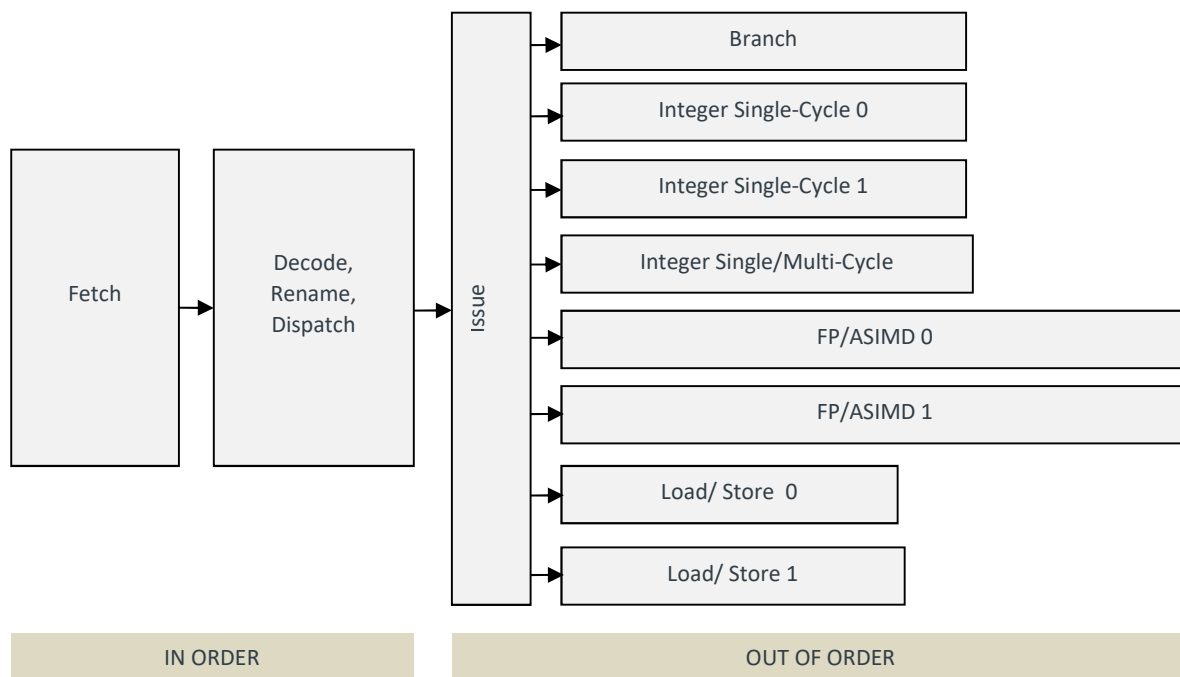
1.3. Scope

This document provides high-level information about the Cortex-A76 pipeline, instruction performance characteristics, and special performance considerations. This information is intended to aid people who are optimizing software and compilers for Cortex-A76. For a more complete description of the Cortex-A76 processor, please refer to the *Cortex-A76 Technical Reference Manual*.

2 Introduction

2.1. Pipeline Overview

The following diagram describes the high-level Cortex-A76 instruction processing pipeline. Instructions are first fetched, then decoded into internal macro-operations (Mops). From there, the Mops proceed through register renaming and dispatch stages. A Mop can be split further into two Uops at dispatch stage. Once dispatched, uops wait for their operands and issue out-of-order to one of eight execution pipelines. Each execution pipeline can accept and complete one uop per cycle.



The execution pipelines support different types of operations, as follows:

Pipeline (mnemonic)	Supported functionality
Branch	Branch μ OPs
Integer Single-Cycle 0/1	Integer ALU μ OPs
Integer Single/Multi-cycle	Integer shift-ALU, multiply, divide, CRC and sum-of-absolute-differences μ OPs
Load/Store Address Generation 0/1	Load, Store and special memory uops
FP/ASIMD-0	ASIMD ALU, ASIMD misc, ASIMD integer multiply, FP convert, FP misc, FP add, FP multiply, FP divide, FP sqrt, crypto uops, store data uops
FP/ASIMD-1	ASIMD ALU, ASIMD misc, FP misc, FP add, FP multiply, ASIMD shift uops, store data uops,

3 Instruction characteristics

3.1. Instruction tables

This chapter describes high-level performance characteristics for most Arm v8.2-A A32, T32 and A64 instructions. A series of tables summarize the effective execution latency and throughput (instruction bandwidth per cycle), pipelines utilized, and special behaviours associated with each group of instructions. Utilized pipelines correspond to the execution pipelines described in chapter 2.

In the tables below, Exec Latency is defined as the minimum latency seen by an operation dependent on an instruction in the described group.

In the tables below, Execution Throughput is defined as the maximum throughput (in instructions per cycle) of the specified instruction group that can be achieved in the entirety of the Cortex-A76 microarchitecture.

3.2. Legend for reading the utilized pipelines

Pipeline name	Symbol used in tables
Branch 0/1	B
Integer single Cycle 0/1	S
Integer single Cycle 0/1 and single/multicycle	I
Integer single/multicycle	M
Integer single Cycle 1 and Integer multicycle	D
Load/Store 0/1	L
FP/ASIMD 0/1	V
FP/ASIMD 0	V0
FP/ASIMD 1	V1

3.3. Branch Instructions

Instruction group	AArch64 instructions	Exec latency	Execution throughput	Utilized Pipelines	Notes
Branch, immed	B	1	1	B	
Branch, register	BR, RET	1	1	B	
Branch and link, immed	BL	1	1	I, B	
Branch and link, register (reg != lr)	BLR	1	1	I, B	
Branch and link, register (reg == lr)	BLR	2	1	I, B	
Compare and branch	CBZ, CBNZ, TBZ, TBNZ	1	1	B	

Instruction group	AArch32 instructions	Exec latency	Execution throughput	Utilized Pipelines	Notes
Branch, immed	B	1	1	B	
Branch, register	BX	1	1	B	
Branch and link, immed	BL, BLX	1	1	I, B	
Branch and link, register (reg != lr)	BLX	1	1	I, B	
Branch and link, register (reg == lr)	BLX	2	1	I, B	
Compare and branch	CBZ, CBNZ	1	1	B	

3.4. Arithmetic and Logical Instructions

Instruction group	AArch64 instructions	Exec latency	Execution throughput	Utilized Pipelines	Notes
Arithmetic, basic	ADD{S}, ADC{S}, SUB{S}, SBC{S}	1	3	I	
Arithmetic, extend and shift	ADD{S}, SUB{S}	2	1	M	
Arithmetic, LSL shift, shift <= 4	ADD{S}, SUB{S}	1	3	I	
Arithmetic, LSR/ASR/ROR shift or LSL shift > 4	ADD{S}, SUB{S}	2	1	M	
Conditional compare	CCMN, CCMP	1	3	I	
Conditional select	CSEL, CSINC, CSINV, CSNEG	1	3	I	
Logical, basic	AND{S}, BIC{S}, EON, EOR, ORN, ORR	1	3	I	
Logical, shift, no flagset	AND, BIC, EON, EOR, ORN, ORR	1	3	I	
Logical, shift, flagset	ANDS, BICS	2	1	M	

Instruction group	AArch32 instructions	Exec latency	Execution throughput	Utilized Pipelines	Notes
ALU, basic	ADD{S}, ADC{S}, ADR, AND{S}, BIC{S}, CMN, CMP, EOR{S}, ORN{S}, ORR{S}, RSB{S}, RSC{S}, SUB{S}, SBC{S}, TEQ, TST	1	3	I	
ALU, shift by register, unconditional	(same as ALU, basic)	2	1	M	
ALU, shift by register, conditional	(same as ALU, basic)	2	1	I, M	
Arithmetic, LSL shift by immed, shift <= 4, unconditional	ADD{S}, ADC{S}, RSB{S}, RSC{S}, SUB{S}, SBC{S}	1	3	I	
Arithmetic, LSL shift by immed, shift <= 4, conditional	ADD{S}, ADC{S}, RSB{S}, RSC{S}, SUB{S}, SBC{S}	1	1	M	

Instruction group	AArch32 instructions	Exec latency	Execution throughput	Utilized Pipelines	Notes
Arithmetic, LSR/ASR/ROR shift by immed or LSL shift by immed > 4	ADD{S}, ADC{S}, RSB{S}, RSC{S}, SUB{S}, SBC{S}	2	1	M	
Logical, shift by immed, noflagset	AND, BIC, EOR, ORN, ORR	1	3	I	
Logical, shift by immed, flagset	AND{S}, BIC{S}, EOR{S}, ORN{S}, ORR{S}	2	1	M	
Test/Compare, shift by immed	CMN, CMP, TEQ, TST	2	1	M	
Branch forms		+1	1	+B	1

Notes:

1. Branch forms are possible when the instruction destination register is the PC. For those cases, an additional branch uop is required. This adds 1 cycle to the latency.

3.5. Move and Shift Instructions

Instruction group	AArch32 instructions	Exec latency	Execution throughput	Utilized Pipelines	Notes
Move, basic	MOV{S}, MOVW, MVN{S}	1	3	I	
Move, shift by immed, no setflags	ASR, LSL, LSR, ROR, RRX, MVN	1	3	I	
Move, shift by immed, setflags	ASRS, LSLS, LSRS, RORS, RRXS, MVNS	2	1	M	
Move, shift by register, no setflags, unconditional	ASR, LSL, LSR, ROR, RRX, MVN	1	3	I	
Move, shift by register, no setflags, conditional	ASR, LSL, LSR, ROR, RRX, MVN	2	3/2	I	
Move, shift by register, setflags	ASRS, LSLS, LSRS, RORS, RRXS, MVNS	2	1	M	
Move, top	MOVT	1	3	I	
Move, branch forms		+1	1	+B	

3.6. Divide and Multiply Instructions

Instruction group	AArch64 instructions	Exec latency	Execution throughput	Utilized Pipelines	Notes
Divide, W-form	SDIV, UDIV	5 to 12	1/12 to 1/5	M	1
Divide, X-form	SDIV, UDIV	5 to 20	1/20 to 1/5	M	1
Multiply accumulate, W-form	MADD, MSUB	2(1)	1	M	2
Multiply accumulate, X-form	MADD, MSUB	4(3)	1/3	M	2, 4
Multiply accumulate long	SMADDL, SMSUBL, UMADDL, UMSUBL	2(1)	1	M	2

Instruction group	AArch64 instructions	Exec latency	Execution throughput	Utilized Pipelines	Notes
Multiply high	SMULH, UMULH	5(3)	1/4	M	5

Instruction group	AArch32 instructions	Exec latency	Execution throughput	Utilized Pipelines	Notes
Divide	SDIV, UDIV	5 to 12	1/12 to 1/5	M	1
Multiply	MUL, SMULBB, SMULBT, SMULTB, SMULTT, SMULWB, SMULWT, SMMUL{R}, SMUAD{X}, SMUSD{X}	2	1	M	
Multiply accumulate, conditional	MLA, MLS, SMLABB, SMLABT, SMLATB, SMLATT, SMLAWB, SMLAWT, SMLAD{X}, SMLS{D}{X}, SMMLA{R}, SMMLS{R}	3	1	M, I	
Multiply accumulate, unconditional	MLA, MLS, SMLABB, SMLABT, SMLATB, SMLATT, SMLAWB, SMLAWT, SMLAD{X}, SMLS{D}{X}, SMMLA{R}, SMMLS{R}	2(1)	1	M	2
Multiply accumulate accumulate long, conditional	UMAAL	4	1	I, M	
Multiply accumulate accumulate long, unconditional	UMAAL	3	1	I, M	
Multiply accumulate long	SMLAL, SMLALBB, SMLALBT, SMLALTB, SMLALTT, SMLALD{X}, SMLS{D}{X}, UMLAL	3	1	M, I	
Multiply long, all setflag, conditional and no setflag	SMULL, UMULL	3	1	M, I	
Multiply long, unconditional and no setflag	SMULL, UMULL	2	1	M	
(Multiply, setflags forms)		+1	(Same as above)	+I	3

Notes:

- Integer divides are performed using a iterative algorithm and block any subsequent divide operations until complete. Early termination is possible, depending upon the data values.
- Multiply-accumulate pipelines support late-forwarding of accumulate operands from similar uops, allowing a typical sequence of multiply-accumulate uops to issue one every N cycles (accumulate latency N shown in parentheses).
- Multiplies that set the condition flags require an additional integer uop.
- X-form multiply accumulates stall the multiplier pipeline for 2 extra cycles.
- Multiply high operations stall the multiplier pipeline for N extra cycles before any other type M uop can be issued to that pipeline, with N shown in parentheses.

3.7. Saturating and Parallel Arithmetic Instructions

Instruction group	AArch32 instructions	Exec latency	Execution throughput	Utilized Pipelines	Notes
Parallel arith, unconditional	SADD16, SADD8, SSUB16, SSUB8, UADD16, UADD8, USUB16, USUB8	2	1	M	
Parallel arith, conditional	SADD16, SADD8, SSUB16, SSUB8, UADD16, UADD8, USUB16, USUB8	2(4)	3/5	M, I	1
Parallel arith with exchange, unconditional	SASX, SSAX, UASX, USAX	3	1	I, M	
Parallel arith with exchange, conditional	SASX, SSAX, UASX, USAX	3(5)	3/5	I, M	1
Parallel halving arith	SHADD16, SHADD8, SHSUB16, SHSUB8, UHADD16, UHADD8, UHSUB16, UHSUB8	2	1	M	
Parallel halving arith with exchange	SHASX, SHSAX, UHASX, UHSAX	3	1	I, M	
Parallel saturating arith	QADD16, QADD8, QSUB16, QSUB8, UQADD16, UQADD8, UQSUB16, UQSUB8	2	2	M	
Parallel saturating arith with exchange	QASX, QSAX, UQASX, UQSAX	3	1	I, M	
Saturate	SSAT, SSAT16, USAT, USAT16	2	1	M	
Saturating arith	QADD, QSUB	2	1	M	
Saturating doubling arith	QDADD, QDSUB	3	1	I, M	

Notes:

1. Conditional GE-setting instructions require three extra uops and two additional cycles to conditionally update the GE field (GE latency shown in parentheses).

3.8. Miscellaneous Data-Processing Instructions

Instruction group	AArch64 instructions	Exec latency	Execution throughput	Utilized Pipelines	Notes
Address generation	ADR, ADRP	1	3	I	
Bitfield extract, one reg	EXTR	1	3	I	
Bitfield extract, two regs	EXTR	3	1	I, M	
Bitfield move, basic	SBFM, UBFM	1	3	I	
Bitfield move, insert	BFM	2	1	M	

Instruction group	AArch64 instructions	Exec latency	Execution throughput	Utilized Pipelines	Notes
Count leading	CLS, CLZ	1	3	I	
Move immed	MOVN, MOVK, MOVZ	1	3	I	
Reverse bits/bytes	RBIT, REV, REV16, REV32	1	3	I	
Variable shift	ASRV, LSLV, LSRV, RORV	1	3	I	

Instruction group	AArch32 instructions	Exec latency	Execution throughput	Utilized Pipelines	Notes
Bit field extract	SBFX, UBFX	1	3	I	
Bit field insert/clear	BFI, BFC	2	1	M	
Count leading zeros	CLZ	1	3	I	
Pack halfword	PKH	2	1	M	
Reverse bits/bytes	RBIT, REV, REV16, REVSH	1	3	I	
Select bytes, unconditional	SEL	1	3	I	
Select bytes, conditional	SEL	2	3/2	I	
Sign/zero extend, normal	SXTB, SXTB, UXTB, UXTB	1	3	I	
Sign/zero extend, parallel	SXTB16, UXTB16	2	1	M	
Sign/zero extend and add, normal	SXTAB, SXTAB, UXTAB, UXTAB	2	1	M	
Sign/zero extend and add, parallel	SXTAB16, UXTAB16	4	1/2	M	
Sum of absolute differences	USAD8, USADA8	2	1	M	

3.9. Load Instructions

The latencies shown assume the memory access hits in the *Level 1* (L1) data cache.

Instruction group	AArch64 instructions	Exec latency	Execution throughput	Utilized Pipelines	Notes
Load register, literal	LDR, LDRSW, PRFM	4	2	L	
Load register, unscaled immed	LDUR, LDURB, LDURH, LDURSB, LDURSH, LDURSW, PRFUM	4	2	L	
Load register, immed post-index	LDR, LDRB, LDRH, LDRSB, LDRSH, LDRSW	4	2	L, I	
Load register, immed pre-index	LDR, LDRB, LDRH, LDRSB, LDRSH, LDRSW	4	2	L, I	

Instruction group	AArch64 instructions	Exec latency	Execution throughput	Utilized Pipelines	Notes
Load register, immed unprivileged	LDTR, LDTRB, LDTRH, LDTRSB, LDTRSH, LDTRSW	4	2	L	
Load register, unsigned immed	LDR, LDRB, LDRH, LDRSB, LDRSH, LDRSW, PRFM	4	2	L	
Load register, register offset, basic	LDR, LDRB, LDRH, LDRSB, LDRSH, LDRSW, PRFM	4	2	L	
Load register, register offset, scale by 4/8	LDR, LDRSW, PRFM	4	2	L	
Load register, register offset, scale by 2	LDRH, LDRSH	5	2	I, L	
Load register, register offset, extend	LDR, LDRB, LDRH, LDRSB, LDRSH, LDRSW, PRFM	4	2	L	
Load register, register offset, extend, scale by 4/8	LDR, LDRSW, PRFM	4	2	L	
Load register, register offset, extend, scale by 2	LDRH, LDRSH	5	2	I, L	
Load pair, signed immed offset, normal, W-form	LDP, LDNP	4	2	L	
Load pair, signed immed offset, normal, X-form	LDP, LDNP	4	1	L	
Load pair, signed immed offset, signed words, base != SP	LDPSW	5	1	I, L	
Load pair, signed immed offset, signed words, base = SP	LDPSW	5	1	I, L	
Load pair, immed post-index, normal	LDP	4	1	L, I	
Load pair, immed post-index, signed words	LDPSW	5	1	I, L	
Load pair, immed pre-index, normal	LDP	4	1	L, I	
Load pair, immed pre-index, signed words	LDPSW	5	1	I, L	

Instruction group	AArch32 instructions	Exec latency	Execution throughput	Utilized Pipelines	Notes
Load, immed offset	LDR{T}, LDRB{T}, LDRD, LDRH{T}, LDRSB{T}, LDRSH{T}	4	2	L	1,2
Load, register offset, plus	LDR, LDRB, LDRD, LDRH, LDRSB, LDRSH	4	2	L	1.2
Load, register offset, minus	LDR, LDRB, LDRD, LDRH, LDRSB, LDRSH	5	2	I, L	1,2
Load, scaled register offset, plus, LSL2	LDR, LDRB	4	2	L	1
Load, scaled register offset, other	LDR, LDRB, LDRH, LDRSB, LDRSH	5	2	I, L	1
Load, immed pre-indexed	LDR, LDRB, LDRD, LDRH, LDRSB, LDRSH	4	2	L, I	1,2
Load, register pre-indexed, shift Rm, plus and minus	LDR, LDRB, LDRH, LDRSB, LDRSH	5	2	I, L, M	3
Load, register pre-indexed	LDRD	4	2	L, I	
Load, register pre-indexed, cond	LDRD	5	1 1/2	L, I	
Load, scaled register pre-indexed, plus, LSL2	LDR, LDRB	4	2	L, I	1
Load, scaled register pre-indexed, unshifted	LDR, LDRB	4	2	L, I	
Load, immed post-indexed	LDR{T}, LDRB{T}, LDRD, LDRH{T}, LDRSB{T}, LDRSH{T}	4	2	L, I	1,2
Load, register post-indexed	LDR, LDRB, LDRH{T}, LDRSB{T}, LDRSH{T}	5	2	I, L	
Load, register post-indexed	LDRD	4	2	L, I	
Load, register post-indexed	LDRT, LDRBT	5	2	I, L	
Load, scaled register post-indexed	LDR, LDRB	4	2	L, M	3
Load, scaled register post-indexed	LDRT, LDRBT	4	2	L, M	3
Preload, immed offset	PLD, PLDW	4	2	L	
Preload, register offset, plus	PLD, PLDW	4	2	L	
Preload, register offset, minus	PLD, PLDW	5	2	I, L	
Preload, scaled register offset, plus LSL2	PLD, PLDW	5	2	I, L	
Preload, scaled register offset, other	PLD, PLDW	5	2	I L	
Load multiple, no writeback, base reg not in list	LDMIA, LDMIB, LDMDA, LDMDB	N	2/R	L	1, 4, 5

Instruction group	AArch32 instructions	Exec latency	Execution throughput	Utilized Pipelines	Notes
Load multiple, no writeback, base reg in list	LDMIA, LDMIB, LDMDA, LDMDB	1+ N	2/R	I, L	1, 4, 5
Load multiple, writeback	LDMIA, LDMIB, LDMDA, LDMDB, POP	1+N	2/R	L, I	1, 4, 5
(Load, all branch forms)		+1		+ B	6

Notes:

1. Condition loads have an extra uop which goes down pipeline I and have 1 cycle extra latency compared to their unconditional counterparts.
2. The throughput of conditional LDRD is 1 as compared to a throughput of 2 for unconditional LDRD.
3. The address update op for addressing forms which use reg scaled reg, or reg extend goes down pipeline 'I' if the shift is LSL where the shift value is less than or equal to 4.
4. N is $\text{floor}[(\text{num_reg}+3)/4]$.
5. R is $\text{floor}[(\text{num_reg}+1)/2]$.
6. Branch forms are possible when the instruction destination register is the PC. For those cases, an additional branch uop is required. This adds 1 cycle to the latency.

3.10. Store Instructions

The following table describes performance characteristics for standard store instructions. Stores uops are split into address and data uops at dispatch time. Once executed, stores are buffered and committed in the background.

Instruction group	AArch64 instructions	Exec latency	Execution throughput	Utilized Pipelines	Notes
Store register, unscaled immed	STUR, STURB, STURH	1	2	L, D	
Store register, immed post-index	STR, STRB, STRH	1	2	L, D	
Store register, immed pre-index	STR, STRB, STRH	1	2	L, D	
Store register, immed unprivileged	STTR, STTRB, STTRH	1	2	L, D	
Store register, unsigned immed	STR, STRB, STRH	1	2	L, D	
Store register, register offset, basic	STR, STRB, STRH	1	2	L, D	
Store register, register offset, scaled by 4/8	STR	1	2	L, D	
Store register, register offset, scaled by 2	STRH	2	3/2	I, L D	
Store register, register offset, extend	STR, STRB, STRH	1	2	L, D	

Instruction group	AArch64 instructions	Exec latency	Execution throughput	Utilized Pipelines	Notes
Store register, register offset, extend, scale by 4/8	STR	1	2	L, D	
Store register, register offset, extend, scale by 1	STRH	2	3/2	I, L, D	
Store pair, immed offset, W-form	STP, STNP	1	2	L, D	
Store pair, immed offset, X-form	STP, STNP	1	1	L, D	
Store pair, immed post-index, W-form	STP	1	1	L, D	
Store pair, immed post-index, X-form	STP	1	1	L, D	
Store pair, immed pre-index, W-form	STP	1	1	L, D	
Store pair, immed pre-index, X-form	STP	1	1	L, D	

Instruction group	AArch32 instructions	Exec latency	Execution throughput	Utilized Pipelines	Notes
Store, immed offset	STR{T}, STRB{T}, STRD, STRH{T}	1	2	L, D	
Store, register offset, plus	STR, STRB, STRD, STRH	1	2	L, D	
Store, register offset, minus	STR, STRB, STRD, STRH	1	2	L, D	
Store, register offset, no shift, plus	STR, STRB	1	2	L, D	
Store, scaled register offset, plus LSL2	STR, STRB	1	2	L, D	
Store, scaled register offset, other	STR, STRB	2	3/2	I, L, D	
Store, scaled register offset, minus	STR, STRB	2	3/2	I, L, D	
Store, immed pre-indexed	STR, STRB, STRD, STRH	1	3/2	I, L, D	
Store, register pre-indexed, plus, no shift	STR, STRB, STRD, STRH	1	3/2	L, D	
Store, register pre-indexed, minus	STR, STRB, STRD, STRH	2	1	I, L, D	
Store, scaled register pre-indexed, plus LSL2	STR, STRB	1	3/2	L, D	
Store, scaled register pre-indexed, other	STR, STRB	2	1	I, L, D, M	1
Store, immed post-indexed	STR{T}, STRB{T}, STRD, STRH{T}	1	3/2	L, D	
Store, register post-indexed	STRH{T}, STRD	1	3/2	L, D	
Store, register post-indexed	STR{T}, STRB{T}	1	3/2	L, D	

Instruction group	AArch32 instructions	Exec latency	Execution throughput	Utilized Pipelines	Notes
Store, scaled register post-indexed	STR{T}, STRB{T}	1	3/2	L, D	
Store multiple, no writeback	STMIA, STMIB, STMDA, STMDB	N	1/N	L, D	2
Store multiple, writeback	STMIA, STMIB, STMDA, STMDB, PUSH	N	1/N	L, D	2

Notes:

1. The address update op for addressing forms which use reg scaled reg, or reg extend goes down pipeline 'I' if the shift is LSL where the shift value is less than or equal to 4.
2. For store multiple instructions, $N = \text{floor}((\text{num_regs} + 3) / 4)$.

3.11. FP Data Processing Instructions

Instruction group	AArch64 instructions	Exec latency	Execution throughput	Utilized Pipelines	Notes
FP absolute value	FABS	2	2	V	
FP arithmetic	FADD, FSUB	2	2	V	
FP compare	FCCMP{E}, FCMP{E}	2	1	V0	
FP divide, H-form	FDIV	7	4/7	V0	1
FP divide, S-form	FDIV	7 to 10	4/9 to 4/7	V0	1
FP divide, D-form	FDIV	7 to 15	1/7 to 2/7	V0	1
FP min/max	FMIN, FMINNM, FMAX, FMAXNM	2	2	V	
FP multiply	FMUL, FNMUL	3	2	V	2
FP multiply accumulate	FMADD, FMSUB, FNMADD, FNMSUB	4 (2)	2	V	3
FP negate	FNEG	2	2	V	
FP round to integral	FRINTA, FRINTI, FRINTM, FRINTN, FRINTP, FRINTX, FRINTZ	3	1	V	
FP select	FCSEL	2	2	V	
FP square root, H-form	FSQRT	7	4/7	V0	1
FP square root, S-form	FSQRT	7 to 10	4/9 to 4/7	V0	1
FP square root, D-form	FSQRT	7 to 17	1/8 to 2/7	V0	1

Instruction group	AArch32 instructions	Exec latency	Execution throughput	Utilized Pipelines	Notes
VFP absolute value	VABS	2	2	V	

Instruction group	AArch32 instructions	Exec latency	Execution throughput	Utilized Pipelines	Notes
VFP arith	VADD, VSUB	2	2	V	
VFP compare, unconditional	VCMP, VCMPE	2	1	V0	
VFP compare, conditional	VCMP, VCMPE	4	1	V, V0	
VFP convert	VCVT{R}, VCVTB, VCVTT, VCVTA, VCVTM, VCVTN, VCVTP	3	1	V0	
VFP divide, H-form	VDIV	7	4/7	V0	1
VFP divide, S-form	VDIV	7 to 10	4/9 to 4/7	V0	1
VFP divide, D-form	VDIV	7 to 15	1/7 to 2/7	V0	1
VFP max/min	VMAXNM, VMINNM	2	2	V	
VFP multiply	VMUL, VNMUL	3	2	V	2
VFP multiply accumulate (chained)	VMLA, VMLS, VNMLA, VNMLS	5 (2)	2	V	3
VFP multiply accumulate (fused)	VFMA, VFMS, VFNMA, VFNMS	4 (2)	2	V	3
VFP negate	VNEG	2	2	V	
VFP round to integral	VRINTA, VRINTM, VRINTN, VRINTP, VRINTR, VRINTX, VRINTZ	3	1	V0	
VFP select	VSELEQ, VSELGE, VSELGT, VSELVS	2	2	V	
VFP square root, H-form	VSQRT	7	4/7	V0	1
VFP square root, S-form	VSQRT	7 to 10	4/9 to 4/7	V0	1
VFP square root, D-form	VSQRT	7 to 17	1/8 to 2/7	V0	1

Notes:

1. FP divide and square root operations are performed using an iterative algorithm and block subsequent similar operations to the same pipeline until complete.
2. FP multiply-accumulate pipelines support late forwarding of the result from FP multiply uops to the accumulate operands of an FP multiply-accumulate uop. The latter can potentially be issued 1 cycle after the FP multiply uop has been issued.
3. FP multiply-accumulate pipelines support late-forwarding of accumulate operands from similar uops, allowing a typical sequence of multiply-accumulate uops to issue one every N cycles (accumulate latency N shown in parentheses).

3.12. FP Miscellaneous Instructions

Instruction group	AArch64 instructions	Exec latency	Execution throughput	Utilized Pipelines	Notes
FP convert, from vec to vec reg	FCVT, FCVTXN	3	1	V0	
FP convert, from gen to vec reg	SCVTF, UCVTF	6	1	M, V0	
FP convert, from vec to gen reg	FCVTAS, FCVTAU, FCVTMS, FCVTMU, FCVTNS, FCVTNU, FCVTPS, FCVTPU, FCVTZS, FCVTZU	4	1	V0, V1	
FP move, immed	FMOV	2	2	V	
FP move, register	FMOV	2	2	V	
FP transfer, from gen to vec reg	FMOV	3	1	M	
FP transfer, from vec to gen reg	FMOV	2	1	V1	

Instruction group	AArch32 instructions	Exec latency	Execution throughput	Utilized Pipelines	Notes
VFP move, immed	VMOV	2	2	V	
VFP move, register	VMOV	2	2	V	
VFP move, insert	VINS	2	2	V	
VFP move, extraction	VMOVX	2	2	V	
VFP transfer, core to vfp, single reg to S-reg, cond	VMOV	5	1	M, V	
VFP transfer, core to vfp, single reg to S-reg, uncond	VMOV	3	1	M	
VFP transfer, core to vfp, single reg to upper/lower half of D-reg	VMOV	5	1	M, V	
VFP transfer, core to vfp, 2 regs to 2 S-reg, cond	VMOV	6	1/2	M, V	
VFP transfer, core to vfp, 2 regs to 2 S-reg, uncond	VMOV	4	1/2	M	
VFP transfer, core to vfp, 2 regs to D-reg, cond	VMOV	5	1	M, V	
VFP transfer, core to vfp, 2 regs to D-reg, uncond	VMOV	3	1	M	
VFP transfer, vfp S-reg or upper/lower half of vfp D-reg to core reg, cond	VMOV	3	1	V1, I	
VFP transfer, vfp S-reg or upper/lower half of vfp D-reg to core reg, uncond	VMOV	2	1	V1	

Instruction group	AArch32 instructions	Exec latency	Execution throughput	Utilized Pipelines	Notes
VFP transfer, vfp 2 S-regs or D-reg to 2 core regs, cond	VMOV	3	1	V1, I	
VFP transfer, vfp 2 S-regs or D-reg to 2 core regs, uncond	VMOV	2	1	V1	

3.13. FP Load Instructions

The latencies shown assume the memory access hits in the Level 1 Data Cache. Compared to standard loads, an extra cycle is required to forward results to FP/ASIMD pipelines.

Instruction group	AArch64 instructions	Exec latency	Execution throughput	Utilized Pipelines	Notes
Load vector reg, literal, S/D/Q forms	LDR		2	L	
Load vector reg, unscaled immed	LDUR	5	2	L	
Load vector reg, immed post-index	LDR	5	2	L, I	
Load vector reg, immed pre-index	LDR	5	2	L, I	
Load vector reg, unsigned immed	LDR	5	2	L, I	
Load vector reg, register offset, basic	LDR	5	2	L, I	
Load vector reg, register offset, scale, S/D-form	LDR	5	2	L, I	
Load vector reg, register offset, scale, H/Q-form	LDR	6	2	I, L	
Load vector reg, register offset, extend	LDR	5	2	L, I	
Load vector reg, register offset, extend, scale, S/D-form	LDR	5	2	L, I	
Load vector reg, register offset, extend, scale, H/Q-form	LDR	6	2	I, L	
Load vector pair, immed offset, S/D-form	LDP, LDNP	5	1	L, I	
Load vector pair, immed offset, Q-form	LDP, LDNP	7	1	L	
Load vector pair, immed post-index, S/D-form	LDP	5	1	I, L	
Load vector pair, immed post-index, Q-form	LDP	7	1	L, I	

Instruction group	AArch64 instructions	Exec latency	Execution throughput	Utilized Pipelines	Notes
Load vector pair, immed pre-index, S/D-form	LDP	5	1	I, L	
Load vector pair, immed pre-index, Q-form	LDP	7	1	L, I	

Instruction group	AArch32 instructions	Exec latency	Execution throughput	Utilized Pipelines	Notes
FP load, register	VLDR	4	2	L	1
FP load multiple, S form	VLDmia, VLDMDB, VPOP	N	2/R	L	1, 2, 3
FP load multiple, D form	VLDmia, VLDMDB, VPOP	N + 2	1/R	L, V	1, 2, 3
(FP load, writeback forms)		(1)		+ I	4

Notes:

1. Condition loads have an extra uop which goes down pipeline V and have 2 cycle extra latency compared to their unconditional counterparts.
2. N is $\text{floor}[(\text{num_reg}+3)/4]$.
3. R is $\text{floor}[(\text{num_reg}+1)/2]$.
4. Writeback forms of load instructions require an extra uop to update the base address. This update is typically performed in parallel with or prior to the load uop (update latency shown in parentheses).

3.14. FP Store Instructions

Stores Mops are split into store address and store data uops at dispatch time. Once executed, stores are buffered and committed in the background.

Instruction group	AArch64 instructions	Exec latency	Execution throughput	Utilized Pipelines	Notes
Store vector reg, unscaled immed, B/H/S/D-form	STUR	2	2	L, I	
Store vector reg, unscaled immed, Q-form	STUR	2	1	L, I	
Store vector reg, immed post-index, B/H/S/D-form	STR	2	2	L, V	
Store vector reg, immed post-index, Q-form	STR	2	1	L, V	
Store vector reg, immed pre-index, B/H/S/D-form	STR	2	2	L, V	
Store vector reg, immed pre-index, Q-form	STR	2	1	L, V	

Instruction group	AArch64 instructions	Exec latency	Execution throughput	Utilized Pipelines	Notes
Store vector reg, unsigned immed, B/H/S/D-form	STR	2	2	L, V	
Store vector reg, unsigned immed, Q-form	STR	2	1	L, V	
Store vector reg, register offset, basic, B/H/S/D-form	STR	2	2	L, V	
Store vector reg, register offset, basic, Q-form	STR	2	1	L, V	
Store vector reg, register offset, scale, H-form	STR	2	2	I, L, V	
Store vector reg, register offset, scale, S/D-form	STR	2	2	L, V	
Store vector reg, register offset, scale, Q-form	STR	2	1	I, L, V	
Store vector reg, register offset, extend, B/H/S/D-form	STR	2	2	L, V	
Store vector reg, register offset, extend, Q-form	STR	2	1	L, V	
Store vector reg, register offset, extend, scale, H-form	STR	2	2	I, L, V	
Store vector reg, register offset, extend, scale, S/D-form	STR	2	2	L, V	
Store vector reg, register offset, extend, scale, Q-form	STR	2	1	I, L, V	
Store vector pair, immed offset, S-form	STP, STNP	2	2	L, V	
Store vector pair, immed offset, D-form	STP, STNP	2	1	L, V	
Store vector pair, immed offset, Q-form	STP, STNP	3	1/2	L, V	
Store vector pair, immed post-index, S-form	STP	2	1	L, V	
Store vector pair, immed post-index, D-form	STP	2	1	L, V	
Store vector pair, immed post-index, Q-form	STP	3	1	L, V	
Store vector pair, immed pre-index, S-form	STP	2	1	L, V	
Store vector pair, immed pre-index, D-form	STP	2	1	L, V	

Instruction group	AArch64 instructions	Exec latency	Execution throughput	Utilized Pipelines	Notes
Store vector pair, immed pre-index, Q-form	STP	3	1/2	L, V	

Instruction group	AArch32 instructions	Exec latency	Execution throughput	Utilized Pipelines	Notes
FP store, immed offset	VSTR	2	2	L I	
FP store multiple, S-form	VSTMIA, VSTMDB, VPUISH	N+1	2/R	L, V	1, 3
FP store multiple, D-form	VSTMIA, VSTMDB, VPUISH	P + 1	1/R	L, V	2, 3
(FP store, writeback forms)		(1)		+ I	4

Notes:

1. For store multiple instructions, $N = \text{floor}((\text{num_regs} + 3) / 4)$.
2. For store multiple instructions, $P = \text{floor}((\text{num_regs} + 1) / 2)$.
3. $R = \text{floor}[(\text{num_regs} + 1) / 2]$.
4. Writeback forms of store instructions require an extra uop to update the base address. This update is typically performed in parallel with or prior to the store uop (update latency shown in parentheses)

3.15. ASIMD Integer Instructions

Instruction group	AArch64 instructions	Exec latency	Execution throughput (T)	Utilized Pipelines	Notes
ASIMD absolute diff	SABD, UABD	2	2	V	
ASIMD absolute diff accum	SABA, UABA	4(1)	1	V1	2
ASIMD absolute diff accum long	SABAL(2), UABAL(2)	4(1)	1	V1	2
ASIMD absolute diff long	SABDL(2), UABDL(2)	2	2	V	
ASIMD arith, basic	ABS, ADD, NEG, SADDL(2), SADDW(2), SHADD, SHSUB, SSUBL(2), SSUBW(2), SUB, UADDL(2), UADDW(2), UHADD, UHSUB, USUBL(2), USUBW(2)	2	2	V	
ASIMD arith, complex	ADDHN(2), RADDHN(2), RSUBHN(2), SQABS, SQADD, SQNEG, SQSUB, SRHADD, SUBHN(2), SUQADD, UQADD, UQSUB, URHADD, USQADD	2	2	V	
ASIMD arith, pair-wise	ADDP, SADDLP, UADDLP	2	2	V	
ASIMD arith, reduce, 4H/4S	ADDV, SADDLV, UADDLV	3	1	V1	

Instruction group	AArch64 instructions	Exec latency	Execution throughput (T)	Utilized Pipelines	Notes
ASIMD arith, reduce, 8B/8H	ADDV, SADDLV, UADDLV	5	1	V1, V	
ASIMD arith, reduce, 16B	ADDV, SADDLV, UADDLV	6	1/2	V1	
ASIMD compare	CMEQ, CMGE, CMGT, CMHI, CMHS, CMLE, CMLT, CMTST	2	1	V	
ASIMD dot product	SDOT, UDOT	2	2	V	
ASIMD logical	AND, BIC, EOR, MOV, MVN, ORN, ORR, NOT	2	2	V	
ASIMD max/min, basic and pair-wise	SMAX, SMAXP, SMIN, SMINP, UMAX, UMAXP, UMIN, UMINP	2	1	V	
ASIMD max/min, reduce, 4H/4S	SMAXV, SMINV, UMAXV, UMINV	3	1	V1	
ASIMD max/min, reduce, 8B/8H	SMAXV, SMINV, UMAXV, UMINV	5	1	V1, V	
ASIMD max/min, reduce, 16B	SMAXV, SMINV, UMAXV, UMINV	6	1/2	V1	
ASIMD multiply, D-form	MUL, SQDMULH, SQRDMULH	4	1	V0	
ASIMD multiply, Q-form	MUL, SQDMULH, SQRDMULH	5	1/2	V0	
ASIMD multiply accumulate, D-form	MLA, MLS	4(1)	1	V0	1
ASIMD multiply accumulate, Q-form	MLA, MLS	5(2)	1/2	V0	1
ASIMD multiply accumulate high, D-form	SQRDMLAH, SQRDMLSH	4	1	V0	-
ASIMD multiply accumulate high, Q-form	SQRDMLAH, SQRDMLSH	5	1/2	V0	-
ASIMD multiply accumulate long	SMLAL(2), SMLSL(2), UMLAL(2), UMLSL(2)	4(1)	1	V0	1
ASIMD multiply accumulate saturating long	SQDMLAL(2), SQDMLSL(2)	4	1	V0	
ASIMD multiply/multiply long (8x8) polynomial, D-form	PMUL, PMULL(2)	3	1	V0	3
ASIMD multiply/multiply long (8x8) polynomial, Q-form	PMUL, PMULL(2)	4	1/2	V0	3
ASIMD multiply long	SMULL(2), UMULL(2), SQDMULL(2)	4	1	V0	
ASIMD pairwise add and accumulate long	SADALP, UADALP	4(1)	1	V1	2
ASIMD shift accumulate	SSRA, SRSRA, USRA, URSRA	4(1)	1	V1	2
ASIMD shift by immed, basic	SHL, SHLL(2), SHRN(2), SSHLL(2), SSHR, SXTL(2), USHLL(2), USHR, UXTL(2)	2	1	V1	

Instruction group	AArch64 instructions	Exec latency	Execution throughput (T)	Utilized Pipelines	Notes
ASIMD shift by immed and insert, basic	SLI, SRI	2	1	V1	
ASIMD shift by immed, complex	RSHRN(2), SQRSHRN(2), SQRSHRUN(2), SQSHL{U}, SQSHRN(2), SQSHRUN(2), SRSHR, UQRSHRN(2), UQSHL, UQSHRN(2), URSHR	4	1	V1	
ASIMD shift by register, basic	SSHL, USHL	2	1	V1	
ASIMD shift by register, complex	SRSHL, SQRSHL, SQSHL, URSHL, UQRSHL, UQSHL	4	1	V1	

Instruction group	AArch32 instructions	Exec latency	Execution throughput (T)	Utilized Pipelines	Notes
ASIMD absolute diff	VABD	2	2	V	
ASIMD absolute diff accum	VABA	4(1)	1	V1	2
ASIMD absolute diff accum long	VABAL	4(1)	1	V1	2
ASIMD absolute diff long	VABDL	2	2	V	
ASIMD arith, basic	VADD, VADDL, VADDW, VNEG, VSUB, VSUBL, VSUBW	2	2	V	
ASIMD arith, complex	VABS, VADDHN, VHADD, VHSUB, VQABS, VQADD, VQNEG, VQSUB, VRADDHN, VRHADD, VRSUBHN, VSUBHN	2	2	V	
ASIMD dot product	VSDOT, VUDOT	2	2	V	
ASIMD arith, pair-wise	VPADD, VPADDL	2	2	V	
ASIMD compare	VCEQ, VCGE, VCGT, VCLE, VTST	2	1	V	
ASIMD logical	VAND, VBIC, VMVN, VORR, VORN, VEOR	2	1	V	
ASIMD max/min	VMAX, VMIN, VPMAX, VPMIN	2	1	V	
ASIMD multiply, D-form	VMUL, VQDMULH, VQRDMULH	4	1	V0	
ASIMD multiply, Q-form	VMUL, VQDMULH, VQRDMULH	5	1/2	V0	
ASIMD multiply accumulate, D-form	VMLA, VMLS	4(1)	1	V0	1
ASIMD multiply accumulate, Q-form	VMLA, VMLS	5(2)	1/2	V0	1
ASIMD multiply accumulate high, D-form	VQRDMLAH, VQRDMLSH	4	1	V0	-

Instruction group	AArch32 instructions	Exec latency	Execution throughput (T)	Utilized Pipelines	Notes
ASIMD multiply accumulate high, Q-form	VQRDMLAH, VQRDMLSH	5	1/2	V0	-
ASIMD multiply accumulate long	VMLAL, VMLSL	4(1)	1	V0	1
ASIMD multiply accumulate saturating long	VQDMLAL, VQDMLSL	4	1	V0	
ASIMD multiply/multiply long (8x8) polynomial, D-form	VMUL (.P8), VMULL (.P8)	3	1	V0	
ASIMD multiply (8x8) polynomial, Q-form	VMUL (.P8)	4	1/2	V0	
ASIMD multiply long	VMULL (.S, .I), VQDMULL	4	1	V0	
ASIMD pairwise add and accumulate	VPADAL	4(1)	1	V1	1
ASIMD shift accumulate	VSRA, VRSRA	4(1)	1	V1	1
ASIMD shift by immed, basic	VMOVL, VSHL, VSHLL, VSHR, VSHRN	2	1	V1	
ASIMD shift by immed and insert, basic	VSLI, VSRI	2	1	V1	
ASIMD shift by immed, complex	VQRSHRN, VQRSHRUN, VQSHL{U}, VQSHRN, VQSHRUN, VRSHR, VRSHRN	4	1	V1	
ASIMD shift by register, basic	VSHL	2	1	V1	
ASIMD shift by register, complex	VQRSHL, VQSHL, VRSHL	4	1	V1	

Notes:

1. Multiply-accumulate pipelines support late-forwarding of accumulate operands from similar uops, allowing a typical sequence of integer multiply-accumulate uops to issue one every cycle or one every other cycle (accumulate latency shown in parentheses).
2. Other accumulate pipelines also support late-forwarding of accumulate operands from similar uops, allowing a typical sequence of such uops to issue one every cycle (accumulate latency shown in parentheses).
3. This category includes instructions of the form “PMULL Vd.8H, Vn.8B, Vm.8B” and “PMULL2 Vd.8H, Vn.16B, Vm.16B”

3.16. ASIMD Floating-Point Instructions

Instruction group	AArch64 instructions	Exec latency	Execution throughput	Utilized Pipelines	Notes
ASIMD FP absolute value/difference	FABS, FABD	2	2	V	
ASIMD FP arith, normal	FABD, FADD, FSUB, FADDP	2	2	V	
ASIMD FP compare	FACGE, FACGT, FCMEQ, FCMGE, FCMGT, FCMLE, FCMLT	2	2	V	
ASIMD FP convert, long (F16 to F32)	FCVTL(2)	4	1/2	V0	
ASIMD FP convert, long (F32 to F64)	FCVTL(2)	3	1	V0	
ASIMD FP convert, narrow (F32 to F16)	FCVTN(2)	4	1/2	V0	
ASIMD FP convert, narrow (F64 to F32)	FCVTN(2), FCVTXN(2)	3	1	V0	
ASIMD FP convert, other, D-form F32 and Q-form F64	FCVTAS, FCVTAU, FCVTMS, FCVTMU, FCVTNS, FCVTNU, FCVTPS, FCVTPU, FCVTZS, FCVTZU, SCVTF, UCVTF	3	1	V0	
ASIMD FP convert, other, D-form F16 and Q-form F32	FCVTAS, VCVTAU, FCVTMS, FCVTMU, FCVTNS, FCVTNU, FCVTPS, FCVTPU, FCVTZS, FCVTZU, SCVTF, UCVTF	4	1/2	V0	
ASIMD FP convert, other, Q-form F16	FCVTAS, VCVTAU, FCVTMS, FCVTMU, FCVTNS, FCVTNU, FCVTPS, FCVTPU, FCVTZS, FCVTZU, SCVTF, UCVTF	6	1/4	V0	
ASIMD FP divide, D-form, F16	FDIV	7	1/7	V0	3
ASIMD FP divide, D-form, F32	FDIV	7 to 10	2/9 to 2/7	V0	3
ASIMD FP divide, Q-form, F16	FDIV	10 to 13	1/13 to 1/10	V0	3
ASIMD FP divide, Q-form, F32	FDIV	7 to 10	1/9 to 1/7	V0	3
ASIMD FP divide, Q-form, F64	FDIV	7 to 15	1/14 to 1/7	V0	3
ASIMD FP max/min, normal	FMAX, FMAXNM, FMIN, FMINNM	2	2	V	
ASIMD FP max/min, pairwise	FMAXP, FMAXNMP, FMINP, FMINNMP	2	2	V	
ASIMD FP max/min, reduce	FMAXV, FMAXNMV, FMINV, FMINNMP	5	2	V	
ASIMD FP max/min, reduce, Q-form F16	FMAXV, FMAXNMV, FMINV, FMINNMP	8	2/3	V	
ASIMD FP multiply	FMUL, FMULX	3	2	V	2

Instruction group	AArch64 instructions	Exec latency	Execution throughput	Utilized Pipelines	Notes
ASIMD FP multiply accumulate	FMLA, FMLS	4(2)	2	V	1
ASIMD FP multiply accumulate long	FMLAL(2), FMLS(2)	5(2)	2	V	1
ASIMD FP negate	FNEG	2	2	V	
ASIMD FP round, D-form F32 and Q-form F64	FRINTA, FRINTI, FRINTM, FRINTN, FRINTP, FRINTX, FRINTZ	3	1	V0	
ASIMD FP round, D-form F16 and Q-form F32	FRINTA, FRINTI, FRINTM, FRINTN, FRINTP, FRINTX, FRINTZ	4	1/2	V0	
ASIMD FP round, Q-form F16	FRINTA, FRINTI, FRINTM, FRINTN, FRINTP, FRINTX, FRINTZ	6	1/4	V0	
ASIMD FP square root, D-form, F16	FSQRT	7	1/7	V0	3
ASIMD FP square root, D-form, F32	FSQRT	7 to 10	2/9 to 2/7	V0	3
ASIMD FP square root, Q-form, F16	FSQRT	11 to 13	1/13 to 1/11	V0	3
ASIMD FP square root, Q-form, F32	FSQRT	7 to 10	1/9 to 1/7	V0	3
ASIMD FP square root, Q-form, F64	FSQRT	7 to 17	1/16 to 1/7	V0	3

Instruction group	AArch32 instructions	Exec latency	Execution throughput	Utilized Pipelines	Notes
ASIMD FP absolute value	VABS	2	2	V	
ASIMD FP arith	VABD, VADD, VPADD, VSUB	2	2	V	
ASIMD FP compare	VACGE, VACGT, VACLE, VACLT, VCEQ, VCGE, VCGT, VCLE	2	2	V	
ASIMD FP convert, integer, D-form	VCVT, VCVTA, VCVTM, VCVTN, VCVTP	3	1	V0	
ASIMD FP convert, integer, Q-form	VCVT, VCVTA, VCVTM, VCVTN, VCVTP	4	1/2	V0	
ASIMD FP convert, fixed, D-form	VCVT	3	1	V0	
ASIMD FP convert, fixed, Q-form	VCVT	4	1/2	V0	
ASIMD FP convert, half-precision	VCVT	4	1/2	V0	
ASIMD FP max/min	VMAX, VMIN, VPMAX, VPMIN, VMAXNM, VMINNM	2	2	V	
ASIMD FP multiply	VMUL, VNMUL	3	2	V	2

Instruction group	AArch32 instructions	Exec latency	Execution throughput	Utilized Pipelines	Notes
ASIMD FP chained multiply accumulate	VMLA, VMLS	5(2)	2	V	1
ASIMD FP fused multiply accumulate	VFMA, VFMS	4(2)	2	V	1
ASIMD FP fused multiply accumulate long	VFMAL(2), VFMSL(2)	4(2)	2	V	1
ASIMD FP negate	VNEG	2	2	V	
ASIMD FP round to integral, D-form	VRINTA, VRINTM, VRINTN, VRINTP, VRINTX, VRINTZ	3	1	V0	
ASIMD FP round to integral, Q-form	VRINTA, VRINTM, VRINTN, VRINTP, VRINTX, VRINTZ	4	1/2	V0	

Notes:

1. ASIMD multiply-accumulate pipelines support late-forwarding of accumulate operands from similar uops, allowing a typical sequence of floating-point multiply-accumulate uops to issue one every N cycles (accumulate latency N shown in parentheses).
2. ASIMD multiply-accumulate pipelines support late forwarding of the result from ASIMD FP multiply uops to the accumulate operands of an ASIMD FP multiply-accumulate uop. The latter can potentially be issued 1 cycle after the ASIMD FP multiply uop has been issued.
3. ASIMD divide and square root operations are performed using an iterative algorithm and block subsequent similar operations to the same pipeline until complete.

3.17. ASIMD Miscellaneous Instructions

Instruction group	AArch64 instructions	Exec latency	Execution throughput	Utilized Pipelines	Notes
ASIMD bit reverse	RBIT	2	2	V	
ASIMD bitwise insert	BIF, BIT, BSL	2	2	V	
ASIMD count	CLS, CLZ, CNT	2	2	V	
ASIMD duplicate, gen reg	DUP	3	1	M	
ASIMD duplicate, element	DUP	2	2	V	
ASIMD extract	EXT	2	2	V	
ASIMD extract narrow	XTN	2	2	V	
ASIMD extract narrow, saturating	SQXTN(2), SQXTUN(2), UQXTN(2)	4	1	V1	
ASIMD insert, element to element	INS	2	2	V	

Instruction group	AArch64 instructions	Exec latency	Execution throughput	Utilized Pipelines	Notes
ASIMD move, FP immed	FMOV	2	2	V	
ASIMD move, integer immed	MOVI, MVNI	2	2	V	
ASIMD reciprocal estimate, D-form F32 and F64	FRECPE, FRECPX, FRSQRTE, URECPE, URSQRTE	3	1	V0	
ASIMD reciprocal estimate, D-form F16 and Q-form F32	FRECPE, FRECPX, FRSQRTE, URECPE, URSQRTE	4	1/2	V0	
ASIMD reciprocal estimate, Q-form F16	FRECPE, FRECPX, FRSQRTE, URECPE, URSQRTE	6	1/4	V0	
ASIMD reciprocal step	FRECPS, FRSQRTS	4	2	V	
ASIMD reverse	REV16, REV32, REV64	2	2	V	
ASIMD table lookup, 1 or 2 table regs	TBL	2	2	V	
ASIMD table lookup, 3 table regs	TBL	4	1/2	V	
ASIMD table lookup, 4 table regs	TBL	4	2/3	V	
ASIMD table lookup extension, 1 table reg	TBX	2	2	V	
ASIMD table lookup extension, 2 table reg	TBX	4	1/2	V	
ASIMD table lookup extension, 3 table reg	TBX	6	2/3	V	
ASIMD table lookup extension, 4 table reg	TBX	6	2/5	V	
ASIMD transfer, element to gen reg	UMOV, SMOV	2	1	V1	
ASIMD transfer, gen reg to element	INS	5	1	M, V	
ASIMD transpose	TRN1, TRN2	2	2	V	
ASIMD unzip/zip	UZP1, UZP2, ZIP1, ZIP2	2	2	V	

Instruction group	AArch32 instructions	Exec latency	Execution throughput	Utilized Pipelines	Notes
ASIMD bitwise insert	VBIF, VBIT, VBSL	2	2	V	
ASIMD count	VCLS, VCLZ, VCNT	2	2	V	
ASIMD duplicate, core reg	VDUP	3	1	M	
ASIMD duplicate, scalar	VDUP	2	2	V	
ASIMD extract	VEXT	2	2	V	
ASIMD move, immed	VMOV, VMVN	2	2	V	

Instruction group	AArch32 instructions	Exec latency	Execution throughput	Utilized Pipelines	Notes
ASIMD move, register	VMOV	2	2	V	
ASIMD move, narrowing	VMOVN	2	2	V	
ASIMD move, saturating	VQMOVN, VQMOVUN	4	1	V1	
ASIMD reciprocal estimate, D-form	VRECPE, VRSQRTE	3	1	V0	
ASIMD reciprocal estimate, Q-form	VRECPE, VRSQRTE	4	1/2	V0	
ASIMD reciprocal step	VRECPS, VRSQRTS	5	2	V	
ASIMD reverse	VREV16, VREV32, VREV64	2	2	V	
ASIMD swap	VSWP	4	2/3	V	
ASIMD table lookup, 1 or 2 table regs	VTBL	2	2	V	
ASIMD table lookup, 3 table regs	VTBL	4	1/2	V	
ASIMD table lookup, 4 table regs	VTBL	4	2/3	V	
ASIMD table lookup extension, 1 reg	VTBX	2	2	V	
ASIMD table lookup extension, 2 table reg	VTBX	4	1/2	V	
ASIMD table lookup extension, 3 table reg	VTBX	6	2/3	V	
ASIMD table lookup extension, 4 table reg	VTBX	8	2/5	V	
ASIMD transfer, scalar to core reg, word	VMOV	2	1	V1	
ASIMD transfer, scalar to core reg, byte/hword	VMOV	3	1	V1, I	
ASIMD transfer, core reg to scalar	VMOV	5	1	M, V	
ASIMD transpose	VTRN	4	2/3	V	
ASIMD unzip/zip	VUZP, VZIP	4	2/3	V	

3.18. ASIMD Load Instructions

The latencies shown assume the memory access hits in the Level 1 Data Cache. Compared to standard loads, an extra cycle is required to forward results to FP/ASIMD pipelines.

Instruction group	AArch64 instructions	Exec latency	Execution throughput	Utilized Pipelines	Notes
ASIMD load, 1 element, multiple, 1 reg, D-form	LD1	5	2	L	
ASIMD load, 1 element, multiple, 1 reg, Q-form	LD1	5	2	L	
ASIMD load, 1 element, multiple, 2 reg, D-form	LD1	5	1	L	
ASIMD load, 1 element, multiple, 2 reg, Q-form	LD1	5	1	L	
ASIMD load, 1 element, multiple, 3 reg, D-form	LD1	6	2/3	L	
ASIMD load, 1 element, multiple, 3 reg, Q-form	LD1	6	2/3	L	
ASIMD load, 1 element, multiple, 4 reg, D-form	LD1	6	1/2	L	
ASIMD load, 1 element, multiple, 4 reg, Q-form	LD1	6	1/2	L	
ASIMD load, 1 element, one lane, B/H/S	LD1	7	2	L, V	
ASIMD load, 1 element, one lane, D	LD1	7	2	L, V	
ASIMD load, 1 element, all lanes, D-form, B/H/S	LD1R	7	2	L, V	
ASIMD load, 1 element, all lanes, D-form, D	LD1R	7	2	L, V	
ASIMD load, 1 element, all lanes, Q-form	LD1R	7	2	L, V	
ASIMD load, 2 element, multiple, D-form, B/H/S	LD2	7	1	L, V	
ASIMD load, 2 element, multiple, Q-form, B/H/S	LD2	7	1	L, V	
ASIMD load, 2 element, multiple, Q-form, D	LD2	7	1	L, V	
ASIMD load, 2 element, one lane, B/H	LD2	7	1	L, V	
ASIMD load, 2 element, one lane, S	LD2	7	1	L, V	

Instruction group	AArch64 instructions	Exec latency	Execution throughput	Utilized Pipelines	Notes
ASIMD load, 2 element, one lane, D	LD2	7	1	L, V	
ASIMD load, 2 element, all lanes, D-form, B/H/S	LD2R	7	1	L, V	
ASIMD load, 2 element, all lanes, D-form, D	LD2R	7	1	L, V	
ASIMD load, 2 element, all lanes, Q-form	LD2R	7	1	L, V	
ASIMD load, 3 element, multiple, D-form, B/H/S	LD3	8	1/2	L, V	
ASIMD load, 3 element, multiple, Q-form, B/H/S	LD3	8	1/2	L, V	
ASIMD load, 3 element, multiple, Q-form, D	LD3	8	1/2	L, V	
ASIMD load, 3 element, one lane, B/H	LD3	7	1/2	L, V	
ASIMD load, 3 element, one lane, S	LD3	7	1/2	L, V	
ASIMD load, 3 element, one lane, D	LD3	7	1/2	L, V	
ASIMD load, 3 element, all lanes, D-form, B/H/S	LD3R	7	1/2	L, V	
ASIMD load, 3 element, all lanes, D-form, D	LD3R	7	1/2	L, V	
ASIMD load, 3 element, all lanes, Q-form, B/H/S	LD3R	7	1/2	L, V	
ASIMD load, 3 element, all lanes, Q-form, D	LD3R	7	1/2	L, V	
ASIMD load, 4 element, multiple, D-form, B/H/S	LD4	8	2/7	L, V	
ASIMD load, 4 element, multiple, Q-form, B/H/S	LD4	10	1/5	L, V	
ASIMD load, 4 element, multiple, Q-form, D	LD4	10	1/5	L, V	
ASIMD load, 4 element, one lane, B/H	LD4	8	1/2	L, V	
ASIMD load, 4 element, one lane, S	LD4	8	1/2	L, V	
ASIMD load, 4 element, one lane, D	LD4	8	1/2	L, V	

Instruction group	AArch64 instructions	Exec latency	Execution throughput	Utilized Pipelines	Notes
ASIMD load, 4 element, all lanes, D-form, B/H/S	LD4R	8	1/2	L, V	
ASIMD load, 4 element, all lanes, D-form, D	LD4R	8	1/2	L, V	
ASIMD load, 4 element, all lanes, Q-form, B/H/S	LD4R	8	1/2	L, V	
ASIMD load, 4 element, all lanes, Q-form, D	LD4R	8	1/2	L, V	
(ASIMD load, writeback form)		(1)		+ I	1

Instruction group	AArch32 instructions	Exec latency	Execution throughput	Utilized Pipelines	Notes
ASIMD load, 1 element, multiple, 1 reg	VLD1	5	2	L	
ASIMD load, 1 element, multiple, 2 reg	VLD1	5	2	L	
ASIMD load, 1 element, multiple, 3 reg	VLD1	5	1	L	
ASIMD load, 1 element, multiple, 4 reg	VLD1	5	1	L	
ASIMD load, 1 element, one lane	VLD1	7	2	L, V	
ASIMD load, 1 element, all lanes, 1 reg	VLD1	7	2	L, V	
ASIMD load, 1 element, all lanes, 2 reg	VLD1	7	2/3	L, V	
ASIMD load, 2 element, multiple, 2 reg	VLD2	7	2/3	L, V	
ASIMD load, 2 element, multiple, 4 reg	VLD2	8	1/2	L, V	
ASIMD load, 2 element, one lane, size 32	VLD2	7	1	L, V	
ASIMD load, 2 element, one lane, size 8/16	VLD2	7	1	L, V	
ASIMD load, 2 element, all lanes	VLD2	7	1	L, V	
ASIMD load, 3 element, multiple, 3 reg	VLD3	8	2/3	L, V	
ASIMD load, 3 element, one lane, size 32	VLD3	8	2/3	L, V	

Instruction group	AArch32 instructions	Exec latency	Execution throughput	Utilized Pipelines	Notes
ASIMD load, 3 element, one lane, size 8/16	VLD3	8	2/3	L, V	
ASIMD load, 3 element, all lanes	VLD3	8	2/3	L, V	
ASIMD load, 4 element, multiple, 4 reg	VLD4	8	1/2	L, V	
ASIMD load, 4 element, one lane, size 32	VLD4	8	1/2	L, V	
ASIMD load, 4 element, one lane, size 8/16	VLD4	8	1/2	L, V	
ASIMD load, 4 element, all lanes	VLD4	8	1/2	L, V	
(ASIMD load, writeback form)		(1)		+I	1

Notes:

1. Writeback forms of load instructions require an extra uop to update the base address. This update is typically performed in parallel with the load uop (update latency shown in parentheses).

3.19. ASIMD Store Instructions

Stores Mops are split into store address and store data uops at dispatch time. Once executed, stores are buffered and committed in the background.

Instruction group	AArch64 instructions	Exec latency	Execution throughput	Utilized Pipelines	Notes
ASIMD store, 1 element, multiple, 1 reg, D-form	ST1	2	2	L, V	
ASIMD store, 1 element, multiple, 1 reg, Q-form	ST1	2	1	L, V	
ASIMD store, 1 element, multiple, 2 reg, D-form	ST1	2	1	L, V	
ASIMD store, 1 element, multiple, 2 reg, Q-form	ST1	3	1/2	L, V	
ASIMD store, 1 element, multiple, 3 reg, D-form	ST1	3	2/3	L, V	
ASIMD store, 1 element, multiple, 3 reg, Q-form	ST1	4	1/3	L, V	
ASIMD store, 1 element, multiple, 4 reg, D-form	ST1	3	1/2	L, V	
ASIMD store, 1 element, multiple, 4 reg, Q-form	ST1	5	1/4	L, V	

Instruction group	AArch64 instructions	Exec latency	Execution throughput	Utilized Pipelines	Notes
ASIMD store, 1 element, one lane, B/H/S	ST1	4	1	V, L	
ASIMD store, 1 element, one lane, D	ST1	4	1	V, L	
ASIMD store, 2 element, multiple, D-form, B/H/S	ST2	4	1	V, L	
ASIMD store, 2 element, multiple, Q-form, B/H/S	ST2	5	1/2	V, L	
ASIMD store, 2 element, multiple, Q-form, D	ST2	5	1/2	V, L	
ASIMD store, 2 element, one lane, B/H/S	ST2	4	1	V, L	
ASIMD store, 2 element, one lane, D	ST2	4	1	V, L	
ASIMD store, 3 element, multiple, D-form, B/H/S	ST3	5	1/2	V, L	
ASIMD store, 3 element, multiple, Q-form, B/H/S	ST3	6	1/3	V, L	
ASIMD store, 3 element, multiple, Q-form, D	ST3	6	1/3	V, L	
ASIMD store, 3 element, one lane, B/H	ST3	4	1/2	V, L	
ASIMD store, 3 element, one lane, S	ST3	4	1/2	V, L	
ASIMD store, 3 element, one lane, D	ST3	5	1/2	V, L	
ASIMD store, 4 element, multiple, D-form, B/H/S	ST4	7	1/3	V, L	
ASIMD store, 4 element, multiple, Q-form, B/H/S	ST4	9	1/6	V, L	
ASIMD store, 4 element, multiple, Q-form, D	ST4	6	1/4	V, L	
ASIMD store, 4 element, one lane, B/H	ST4	5		V, L	
ASIMD store, 4 element, one lane, S	ST4		2/3	V, L	
ASIMD store, 4 element, one lane, D	ST4			V, L	
(ASIMD store, writeback form)		(1)		Add I	1

Instruction group	AArch32 instructions	Exec latency	Execution throughput	Utilized Pipelines	Notes
ASIMD store, 1 element, multiple, 1 reg	VST1	2	2	L, V	
ASIMD store, 1 element, multiple, 2 reg	VST1	2	2	L, V	
ASIMD store, 1 element, multiple, 3 reg	VST1	3	2/3	L, V	
ASIMD store, 1 element, multiple, 4 reg	VST1	3	1/2	L, V	
ASIMD store, 1 element, one lane	VST1	4	2	V, L	
ASIMD store, 2 element, multiple, 2 reg	VST2	4	1	V, L	
ASIMD store, 2 element, multiple, 4 reg	VST2	5	1/2	V, L	
ASIMD store, 2 element, one lane	VST2	4	2	V, L	
ASIMD store, 3 element, multiple, 3 reg	VST3	5	2/3	V, L	
ASIMD store, 3 element, one lane, size 32	VST3	4	1	V, L	
ASIMD store, 3 element, one lane, size 8/16	VST3	4	1	V, L	
ASIMD store, 4 element, multiple, 4 reg	VST4	8	1/2	V, L	
ASIMD store, 4 element, one lane, size 32	VST4	7	2	V, L	
ASIMD store, 4 element, one lane, size 8/16	VST4	7	2	V, L	
(ASIMD store, writeback form)		(1)		+I	1

Notes:

1. Writeback forms of store instructions require an extra uop to update the base address. This update is typically performed in parallel with the store uop (update latency shown in parentheses).

3.20. Cryptography Extensions

Instruction group	AArch64 instructions	Exec latency	Execution throughput	Utilized Pipelines	Notes
Crypto AES ops	AESD, AESE, AESIMC, AESMC	2	1	V0	
Crypto polynomial (64x64) multiply long	PMULL(2)	2	1	V0	
Crypto SHA1 hash acceleration op	SHA1H	2	1	V0	
Crypto SHA1 hash acceleration ops	SHA1C, SHA1M, SHA1P	4	1	V0	
Crypto SHA1 schedule acceleration ops	SHA1SU0, SHA1SU1	2	1	V0	
Crypto SHA256 hash acceleration ops	SHA256H, SHA256H2	4	1	V0	
Crypto SHA256 schedule acceleration ops	SHA256SU0, SHA256SU1	2	1	V0	

Instruction group	AArch32 instructions	Exec latency	Execution throughput	Utilized Pipelines	Notes
Crypto AES ops	AESD, AESE, AESIMC, AESMC	2	1	V0	1
Crypto polynomial (64x64) multiply long	VMULL.P64	2	1	V0	
Crypto SHA1 hash acceleration op	SHA1H	2	1	V0	
Crypto SHA1 hash acceleration ops	SHA1C, SHA1M, SHA1P	4	1	V0	
Crypto SHA1 schedule acceleration ops	SHA1SU0, SHA1SU1	2	1	V0	
Crypto SHA256 hash acceleration ops	SHA256H, SHA256H2	4	1	V0	
Crypto SHA256 schedule acceleration ops	SHA256SU0, SHA256SU1	2	1	V0	

Notes:

1. Adjacent AESE/AESMC instruction pairs and adjacent AESD/AESIMC instruction pairs will exhibit the performance characteristics described in Section 4.5.

3.21. CRC

Instruction group	AArch64 instructions	Exec latency	Execution throughput	Utilized Pipelines	Notes
CRC checksum ops	CRC32, CRC32C	2	1	M	1

Instruction group	AArch32 instructions	Exec latency	Execution throughput	Utilized Pipelines	Notes
CRC checksum ops	CRC32, CRC32C	2	1	M	1

Notes:

1. CRC execution supports late forwarding of the result from a producer uop to a consumer uop. This results in a 1 cycle reduction in latency as seen by the consumer.

4 Special considerations

4.1. Dispatch Constraints

Dispatch of uops from the in-order portion to the out-of-order portion of the microarchitecture includes a number of constraints. It is important to consider these constraints during code generation in order to maximize the effective dispatch bandwidth and subsequent execution bandwidth of Cortex-A76.

The dispatch stage can process up to 4 Mops per cycle and dispatch up to 8 uops per cycle, with the following limitations on the number of uops of each type that may be simultaneously dispatched.

- Up to 2 uops utilizing B pipeline
- Up to 4 uops utilizing S pipelines
- Up to 2 uops utilizing M pipeline
- Up to 2 uops utilizing each of the V pipelines.
- Up to 2 uops utilizing each of the L pipelines

In the event there are more uops available to be dispatched in a given cycle than can be supported by the constraints above, uops will be dispatched in oldest to youngest age-order to the extent allowed by the above.

4.2. Dispatch Stall

In the event of a V-pipeline uop containing more than 1 quad-word register source, a portion or all of which was previously written as one or multiple single words, that uop will stall in dispatch for three cycles. This stall occurs only on the first such instance, and subsequent consumers of the same register will not experience this stall.

4.3. Optimizing General-Purpose Register Spills and Fills

Register transfers between general-purpose registers (GPR) and ASIMD registers (VPR) are lower latency than reads and writes to the cache hierarchy, thus it is recommended that GPR registers be filled/spilled to the VPR rather to memory, when possible.

4.4. Optimizing Memory Copy

The Cortex-A76 processor includes two load/store pipelines, which allow it to execute two 128 bit load uops and one 128 bit store uop every cycle

To achieve maximum throughput for memory copy (or similar loops), one should do the following.

- Unroll the loop to include multiple load and store operations per iteration, minimizing the overheads of looping.
- Use discrete, non-writeback forms of load and store instructions while interleaving them.
- Align stores on 16B boundary wherever possible.
- The following example shows a recommended instruction sequence for a long memory copy in AArch64 state:

Loop_start:

```
SUBS    X2,X2,#192
LDP     Q3,Q4,[x1,#0]
STP     Q3,Q4,[x0,#0]
LDP     Q3,Q4,[x1,#32]
STP     Q3,Q4,[x0,#32]
```

```

LDP    Q3,Q4,[x1,#64]
STP    Q3,Q4,[x0,#64]
LDP    Q3,Q4,[x1,#96]
STP    Q3,Q4,[x0,#96]
LDP    Q3,Q4,[x1,#128]
STP    Q3,Q4,[x0,#128]
LDP    Q3,Q4,[x1,#160]
STP    Q3,Q4,[x0,#160]
ADD    X1,X1,#192
ADD    X0,X0,#192
BGT    Loop_start

```

A recommended copy routine for AArch32 would look similar to the sequence above but would use LDRD/STRD instructions. Avoid load-/store-multiple instruction encodings (such as LDM and STM).

4.5. Load/Store Alignment

The Arm v8.2-A architecture allows many types of load and store accesses to be arbitrarily aligned. The Cortex-A76 processor handles most unaligned accesses without performance penalties. However, there are cases which reduce bandwidth or incur additional latency, as described below.

- Load operations that cross a cache-line (64-byte) boundary.
- Quad-word load operations that are not 4B aligned.
- Store operations that cross a 16B boundary.

4.6. AES Encryption/Decryption

Cortex-A76 can issue one AESE/AESMC/AESD/AESIMC instruction every cycle (fully pipelined) with an execution latency of two cycles. This means encryption or decryption for at least two data chunks should be interleaved for maximum performance:

```

AESE   data0, key0
AESMC  data0, data0
AESE   data1, key0
AESMC  data1, data1
AESE   data0, key0
AESMC  data0, data0
AESE   data1, key1
AESMC  data1, data1
...

```

Pairs of dependent AESE/AESMC and AESD/AESIMC instructions are higher performance when they are adjacent in the program code and both instructions use the same destination register.

4.7. Region Based Fast Forwarding

The forwarding logic in the V pipelines is optimized to provide optimal latency for instructions which are expected to commonly forward to one another. These optimized forwarding regions are defined as follows:

Region	Instruction Types
1	ASIMD ALU, ASIMD shift and certain ASIMD Miscellaneous.
2	FP multiply, FP multiply-accumulate, FP compare, FP add/sub and certain ASIMD Miscellaneous.
3	Cryptography Extensions.

In addition to the regions mentioned in the table above, all floating point and ASIMD instructions can fast forward to FP and ASIMD stores

Exceptions to these forwarding regions are as follows:

- Fast forwarding will not occur in AArch32 mode if the consuming register's width is greater than that of the producer.
- Element sources used by FP multiply and multiply-accumulate operations cannot be consumers.
- Complex ASIMD shift by immediate/register and shift accumulate instructions cannot be producers (see section 3.14) in region 1.
- ASIMD extract narrow, saturating instructions cannot be producers (see section 3.16) in region 1.
- ASIMD absolute difference accumulate and pairwise add and accumulate instructions cannot be producers (see section 3.14) in region 1.
- For FP producer-consumer pairs, the precision of the instructions should match (single, double or half) in region 2.
- Pair-wise FP instructions cannot be producers or consumers in region 2.

The effective latency of FP and ASIMD instructions as described in section 3 is increased by one cycle if the producer and consumer instructions are not part of the same forwarding region, or if they are affected by the exceptions described above. ASIMD integer multiply/multiply-accumulate, ASIMD max/min reduction, FP divide and square root, FP convert/round and reciprocal/reciprocal sqrt estimate instructions are not subject to this behavior.

4.8. Branch instruction alignment

Branch instruction and branch target instruction alignment and density can affect performance. For best-case performance, consider the following guidelines.

- Avoid placing more than 4 branch instructions within an aligned 32-byte instruction memory region.
- When possible, a branch and its target should be located within the same 2M aligned memory region.

Consider aligning subroutine entry points and branch targets to 32B boundaries, within the bounds of the code-density requirements of the program. This will ensure that the subsequent fetch can maximize bandwidth following the taken branch by bringing in all useful instructions.

For loops which comprise 32 or fewer instruction bytes, it is preferred that the loop be located entirely within a single aligned 32-byte instruction memory region.

4.9. FPCR self-synchronization

Programmers and compiler writers should note that writes to the FPCR register are self-synchronizing, i.e. its effect on subsequent instructions can be relied upon without an intervening context synchronizing operation.

4.10. Special Register Access

The Cortex-A76 processor performs register renaming for general purpose registers to enable speculative and out-of-order instruction execution. But most special-purpose registers are not renamed. Instructions that read or write non-renamed registers are subjected to one or more of the following additional execution constraints.

- Non-Speculative Execution – Instructions may only execute non-speculatively.
- In-Order Execution – Instructions must execute in-order with respect to other similar instructions or in some cases all instructions.
- Flush Side-Effects – Instructions trigger a flush side-effect after executing for synchronization.

The table below summarizes various special-purpose register read accesses and the associated execution constraints or side-effects.

Register Read	Non-Speculative	In-Order	Flush Side-Effect	Notes
APSR	Yes	Yes	No	3
CurrentEL	No	Yes	No	
DAIF	No	Yes	No	
DLR_ELO	No	Yes	No	
DSPSR_ELO	No	Yes	No	
ELR_*	No	Yes	No	
FPCR	No	Yes	No	
FPSCR	Yes	Yes	No	2
FPSR	Yes	Yes	No	2
NZCV	No	No	No	1
SP_*	No	No	No	1
SPSel	No	Yes	No	
SPSR_*	No	Yes	No	

Notes:

1. The NZCV and SP registers are fully renamed.
2. FPSR/FPSCR reads must wait for all prior instructions that may update the status flags to execute and retire.
3. APSR reads must wait for all prior instructions that may set the Q bit to execute and retire.

The table below summarizes various special-purpose register write accesses and the associated execution constraints or side-effects.

Register Write	Non-Speculative	In-Order	Flush Side-Effect	Notes
APSR	Yes	Yes	No	4
DAIF	Yes	Yes	No	
DLR_ELO	Yes	Yes	No	
DSPSR_ELO	Yes	Yes	No	
ELR_*	Yes	Yes	No	

Register Write	Non-Speculative	In-Order	Flush Side-Effect	Notes
FPCR	Yes	Yes	Maybe	2
FPSCR	Yes	Yes	Maybe	2, 3
FPSR	Yes	Yes	No	3
NZCV	No	No	No	1
SP_*	No	No	No	1
SPSel	Yes	Yes	Yes	
SPSR_*	Yes	Yes	No	

Notes:

1. The NZCV and SP registers are fully renamed.
2. If the FPCR/FPSCR write is predicted to change the control field values, it will introduce a barrier which prevents subsequent instructions from executing. If the FPCR/FPSCR write is predicted to not change the control field values, it will execute without a barrier but trigger a flush if the values change.
3. FPSR/FPSCR writes must stall at dispatch if another FPSR/FPSCR write is still pending.
4. APSR writes that set the Q bit will introduce a barrier which prevents subsequent instructions from executing until the write completes.

4.11. Register Forwarding Hazards

The Armv8-A architecture allows FP/ASIMD instructions to read and write 32-bit S-registers. In AArch32, Each S-register corresponds to one half (upper or lower) of an overlaid 64-bit D-register. A Q register in turn consists of two overlaid D register. Register forwarding hazards may occur when one uop reads a Q-register operand that has recently been written with one or more S-register result. Consider the following scenario.

```
VADD    S0, S1, S2
```

```
VADD    Q6, Q5, Q0
```

The first instruction writes S0, which correspond to the lowest part of Q0. The second instruction then requires Q0 as an input operand. In this scenario, there is a dependency RAW dependency between the first and the second instructions. In most cases, Cortex-A76 performs slightly worse in such situations.

Cortex-A76 is able to avoid this register-hazard condition for certain cases. The following rules describe the conditions under which a register-hazard can occur.

- The producer writes an S-register (not a D[x] scalar)
- The consumer reads an overlapping Q-register (not as a D[x] scalar)
- The consumer is a FP/ASIMD uop (not a store or MOV uop)

To avoid unnecessary hazards, it is recommended that the programmer use D[x] scalar writes when populating registers prior to ASIMD operations. For example, either of the following instruction forms would safely prevent a subsequent hazard.

```
VLD1.32    D0[x], [address]
```

```
VADD        Q1, Q0, Q2F
```

4.12. IT Blocks

The Armv8-A architecture performance deprecates some uses of the IT instruction in such a way that software may be written using multiple naïve single instruction IT blocks. It is preferred that software instead generate multi instruction IT blocks rather than single instruction blocks.